

Lecture #4

# GPU Rendering

Computer Graphics  
Winter Term 2016/17

Marc Stamminger

# GPUs

- Rendering is very compute intensive
  - Full-HD, 60 frames per second → 120 mio pixels to be set per second !
  - Today's scene have easily 10-100 millions of triangles
- But it can also be easily parallelized
  - per triangle, per vertex, per pixel, ...

- GPUs: Graphics Processing Units

- a processor like a CPU
- but with many cores (hundreds or thousands)
- and with its own memory

- APIs:

- OpenGL, DirectX, ...
- command-based
- In this lecture: OpenGL / WebGL




NVIDIA Titan GTX  
3.000 cores

# GPUs

- Data Transfer

- CPU – Main Memory: fast
- GPU – GPU Memory: very fast
- CPU – GPU: slow



- **We first have to submit the scene data to GPU memory**
- **And then we can start the render process**

# Minimal WebGL program

```
// get WebGL context  
var gl = canvas.getContext("webgl");
```

```
// triangle coordinates  
var v = [-1,-1, 1,-1, 0,0];
```

1

```
-> upload v to GPU Memory  
gl.???
```

2

```
-> setup shaders  
gl.???
```

3

```
-> render  
gl.???
```


4

# Triangle Data 1

- One single triangle:

```
var v = [-1,-1, 1,-1, 0,1];
```


→ OpenGL coordinates go from -1 to 1 ! (for now)



- Two triangles:

```
var v = [-1,0, 1,0, 0,1, -1,0, 1,0, 0,-1];
```

→ inefficient, because two vertices are used twice




- Indexed Face Set data structure:

```
var v = [-1,0, 1,0, 0,1, 0,-1];  
var i = [0,1,2, 0,3,1];
```

vertex  
coordinates

vertex indices  
per triangle

- → Needed for large scenes with many triangles (millions)



# Upload Data to GPU Memory

2


- Upload vertex array to ARRAY\_BUFFER

```
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);
```

- memory chunks on the GPU are referred to as “buffer objects”
- `gl.createBuffer()` creates such an object
- one such buffer object contains the current vertex array data
- setting this `ARRAY_BUFFER` is called “binding the buffer”: `gl.bindBuffer(...)`
- data can be copied into the currently bound buffer using `gl.bufferData(...)`
- JavaScript arrays first have to be converted to a `TypedArray`, in this example `Float32Array`

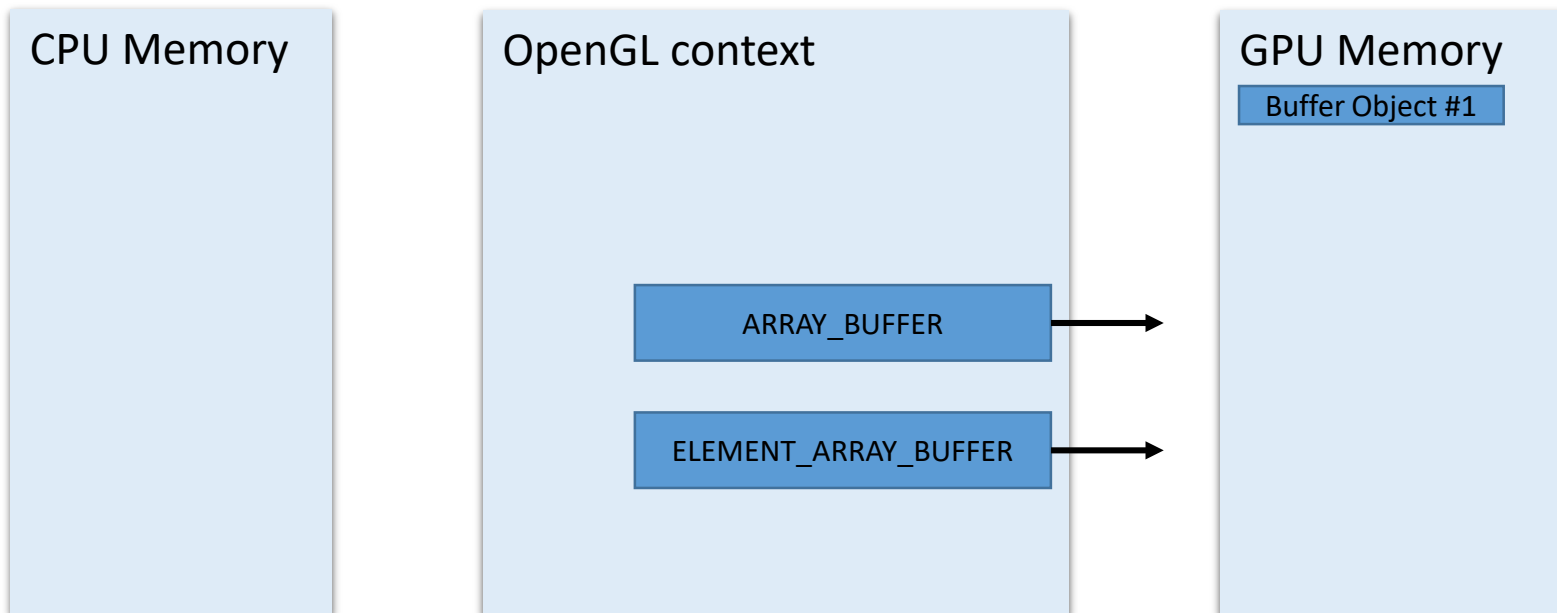
# Upload Data to GPU Memory

2




```
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);
```

- `createBuffer()` creates a GPU memory object
- first, this is only a handle, no memory is allocated yet



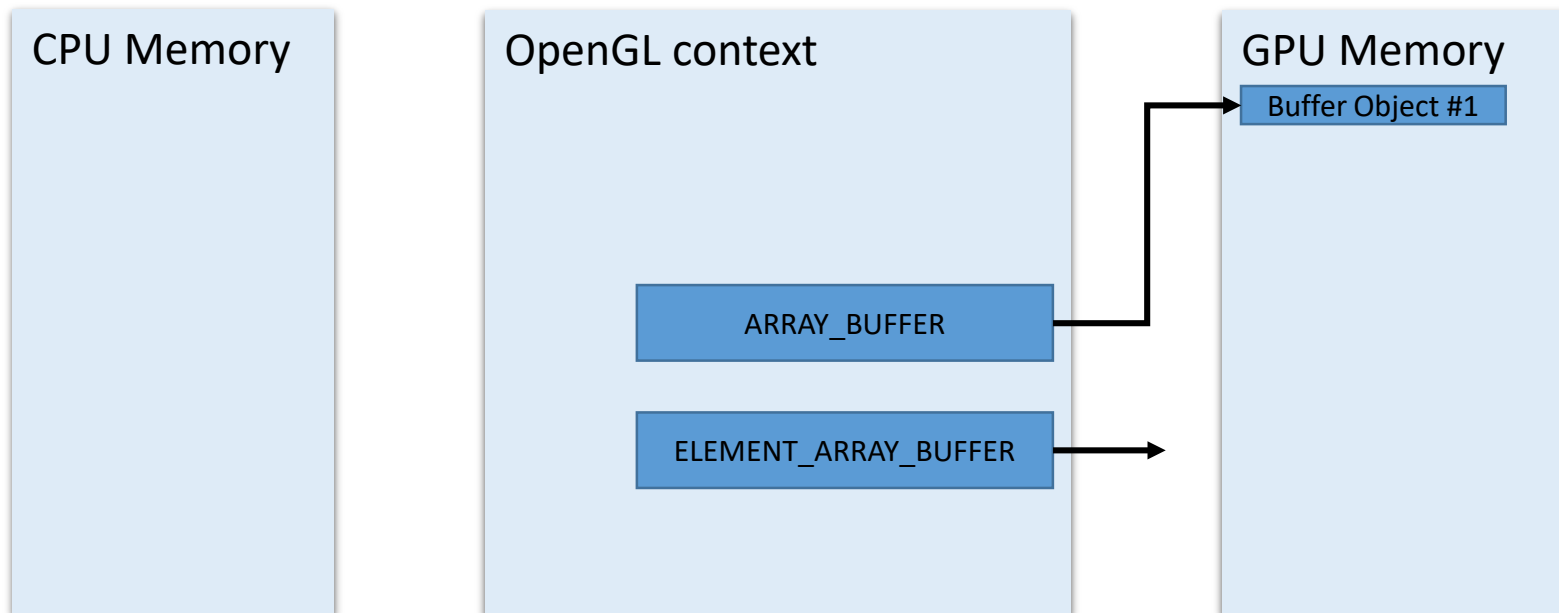
# Upload Data to GPU Memory

2



```
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);
```


- `bindBuffer(...)` connects this buffer with OpenGL's **array buffer**
- later render commands get their data from the currently bound array buffer





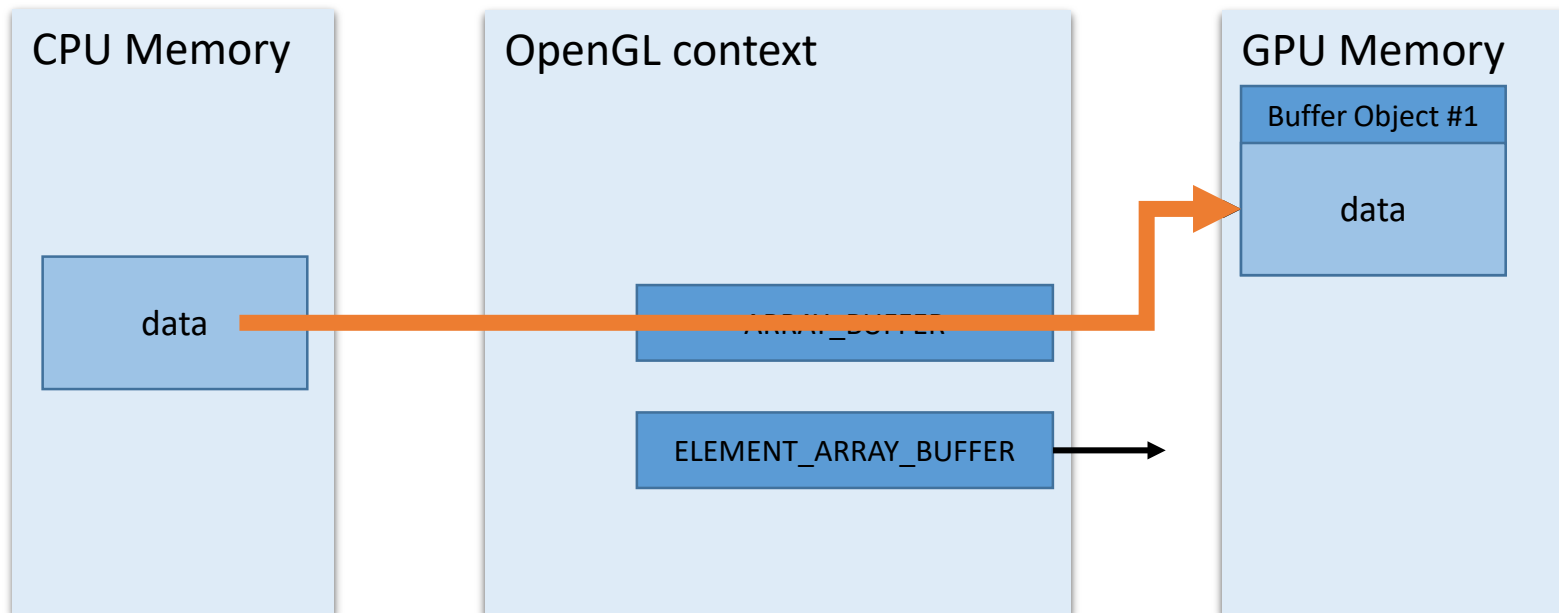
# Upload Data to GPU Memory

2



```
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);
```

- `bufferData(...)` loads data to GPU memory
- in this step, memory is allocated in GPU memory



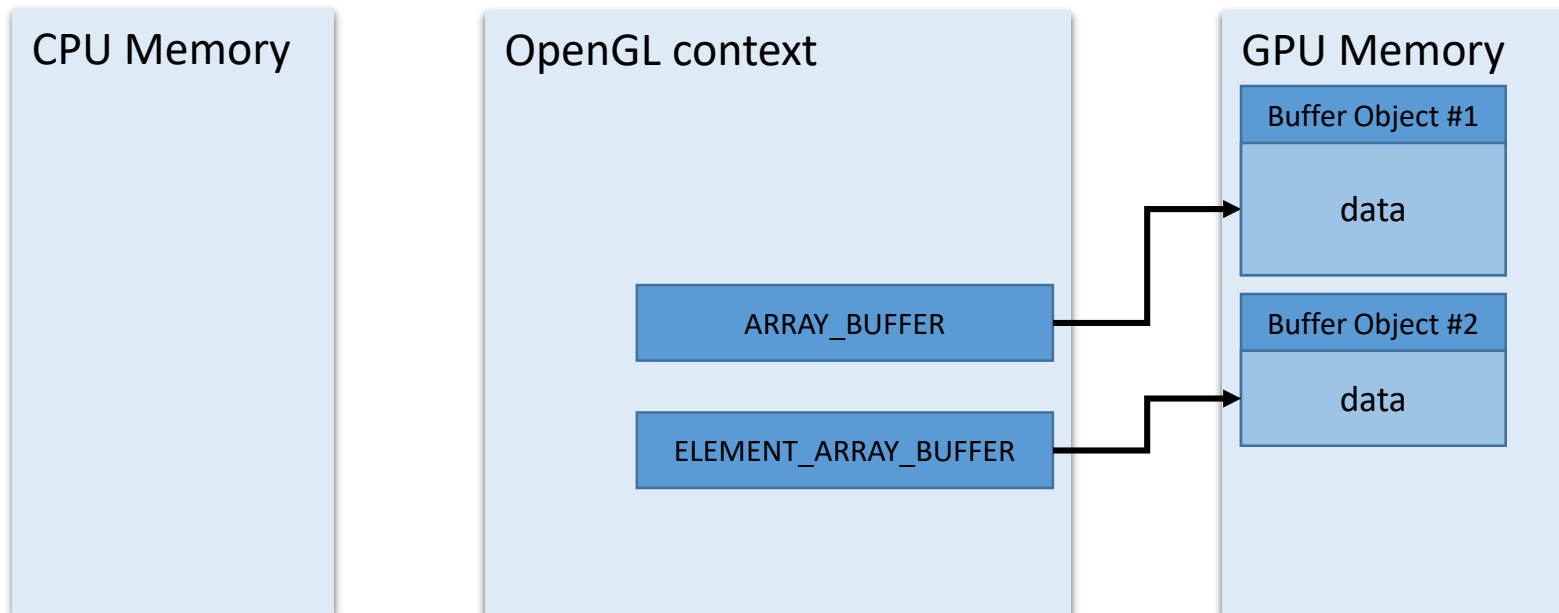
# Upload Data to GPU Memory

2

- We do the same with the `ELEMENT_ARRAY_BUFFER`

```
var ibo = gl.createBuffer();  
gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, ibo);  
gl.bufferData(gl.ELEMENT_ARRAY_BUFFER, new Uint16Array(i), gl.STATIC_DRAW);
```

- index data needs to be bound to `ELEMENT_ARRAY_BUFFER`
- index data must be an unsigned integer array



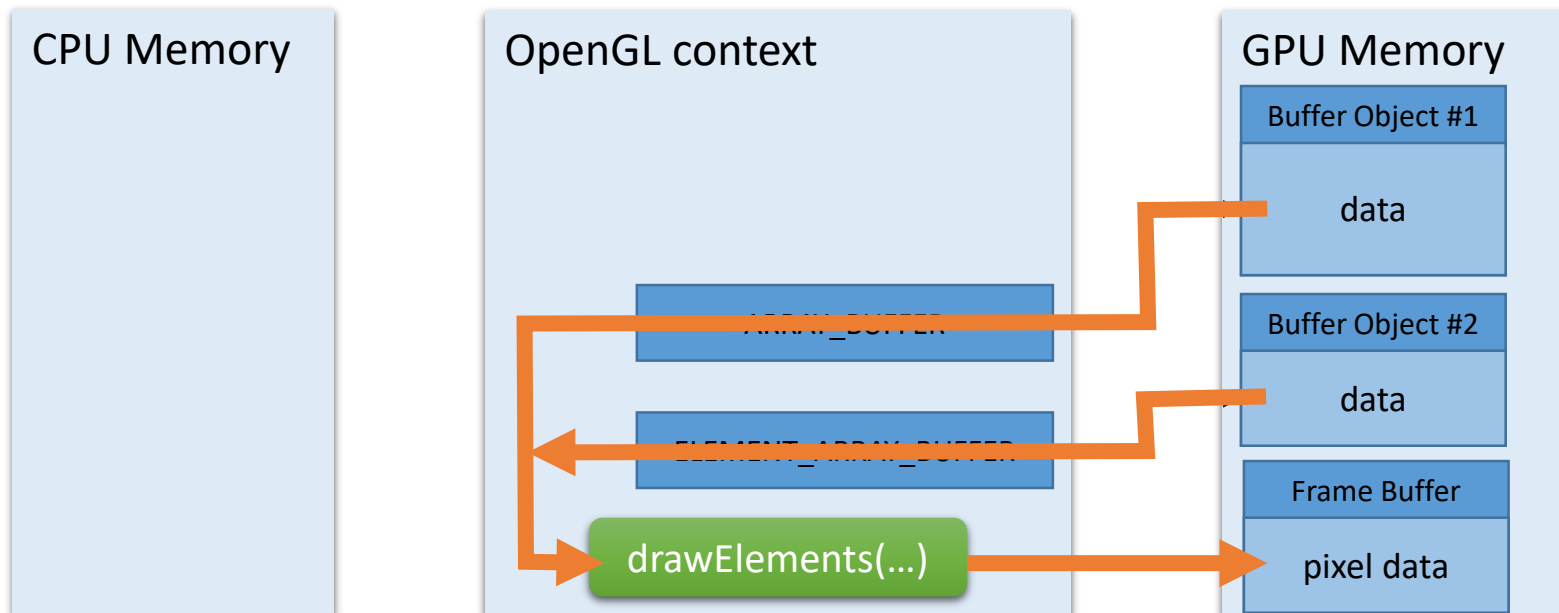
# Upload Data to GPU Memory

2

- Finally, we will render the triangles

```
gl.drawElements(...)
```

- Vertex data will be fetched from the ARRAY\_BUFFER
- and index data from the ELEMENT\_ARRAY\_BUFFER



# Setup Shaders

3

- Last, we have to define a **Vertex Shader** and a **Fragment Shader**
- A shader is a little program written in a C-like language **GLSL**
- The vertex shader is called for every vertex before rasterization
  - here we can transform the vertices (e.g. to animate objects, to zoom, ...)  
→ next lecture “Transformations”
  - and we can add attributes to the vertices that are interpolated during rasterization  
→ Gouraud Shading from previous lecture
- The fragment shader is called for every rasterized pixel
  - here we compute the color of a pixel  
→ later lectures “Lighting”, “Texturing”, ...

# Setup Shaders

3

- Vertex Shader

```
attribute vec2 pos;  
  
void main(void) {  
    gl_Position = vec4(pos, 0.0, 1.0);  
}
```

- executed for every vertex

```
var v = [-1,0, 1,0, 0,1, -1,0, 1,0, 0,-1];
```



pos



vertex shader



gl\_Position



rasterizer

# Setup Shaders

3

- Vertex Shader

```
attribute vec2 pos;

void main(void) {
    gl_Position = vec4(pos, 0.0, 1.0);
}
```

- Bind attributes from ARRAY\_BUFFER to vertex shader attributes:

```
var attrPos = gl.getAttribLocation(shaderProgram, "pos");
gl.enableVertexAttribArray(attrPos);
gl.vertexAttribPointer(attrPos, 2, gl.FLOAT, false, 8, 0);
```

2D-attribute

Offset between  
successive vertices

Start offset of  
first vertex

# Setup Shaders


3

- Interleaved attributes

vertex 0

attribute 1 "pos": (x0,y0)

attribute 2 "col": (r0,g0,b0)



ARRAY\_BUFFER

20 bytes



offset  
"pos"  
= 0

offset  
"col"  
= 8

```
var attrPos = gl.getAttributeLocation(shaderProgram, "pos");  
gl.enableVertexAttribArray(attrPos);  
gl.vertexAttribPointer(attrPos, 2, gl.FLOAT, false, 20, 0);
```

```
var attrCol = gl.getAttributeLocation(shaderProgram, "col");  
gl.enableVertexAttribArray(attrCol);  
gl.vertexAttribPointer(attrCol, 3, gl.FLOAT, false, 20, 8);
```

# Setup Shaders

3

- A more complicated shader

```
uniform vec2 offset;  
uniform float zoom;  
attribute vec2 pos;  
  
void main(void) {  
    gl_Position = vec4((pos+offset)*zoom, 0.0, 1.0);  
}
```


- **uniform**: variables that can be set globally
- **vec2, vec3, vec4, mat2, mat3, mat4**:
  - special types for graphics computing
  - includes many functions, e.g. for dot product, matrix multiplication etc.
  - see [documentation of GLSL](#)



# Setup Shaders

3

- Next, the triangles is rasterized



- for every pixel, a fragment shader is called (aka pixel shader):

```
precision highp float; // set precision of float computations

void main(void){
    gl_FragColor = vec4(1,0,0,1);
}
```

- finally, pixel is set to gl\_FragColor

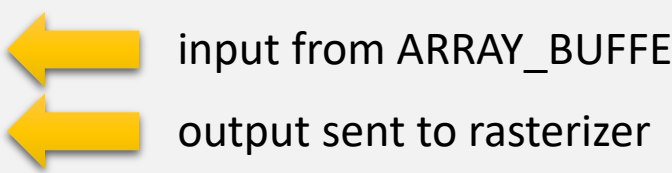
# Setup Shaders

3

- vertex shaders can output additional values, which are then interpolated during rasterization (→ Gouraud Shading)
- these become an input of the fragment shader:

```
// vertex shader
attribute vec2 pos;
attribute vec3 col;
varying vec3 c;

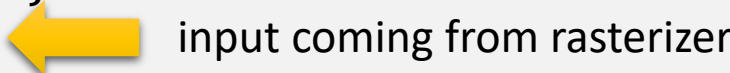
void main(void) {
    gl_Position = vec4((pos+offset)*zoom, 0.0, 1.0);
    c = col;
}
```



Two yellow arrows point from the text to the code. The first arrow points from "input from ARRAY\_BUFFER" to the `attribute` lines (`pos` and `col`). The second arrow points from "output sent to rasterizer" to the `varying` line (`c`).

```
// fragment shader
precision highp float;
varying vec3 c;

void main(void) {
    gl_FragColor = vec4(c,1);
}
```



A yellow arrow points from the text "input coming from rasterizer" to the `varying` line (`c`).

# Setup Shaders

3

- Vertex shader and Fragment shaders come in pairs, interacting with varyings
- Shaders are compiled using GL commands:

```
var vertexShader = gl.createShader(gl.VERTEX_SHADER);  
gl.shaderSource(vertexShader, "shaderCodeAsString");  
gl.compileShader(vertexShader);  
...  
var fragmentShader = gl.createShader(gl.FRAGMENT_SHADER);  
gl.shaderSource(fragmentShader, "shaderCodeAsString");  
gl.compileShader(fragmentShader);
```

- and must then be linked to a **shader program**:

```
var shaderProgram = gl.createProgram();  
gl.attachShader(shaderProgram, vertexShader);  
gl.attachShader(shaderProgram, fragmentShader);  
gl.linkProgram(shaderProgram);
```

- and finally be activated:

```
gl.useProgram(shaderProgram);
```

# Render

4

- Finally, we can make a render call:

```
gl.drawElements(GL_TRIANGLES, 6, GL_UNSIGNED_SHORT, 0);
```

process 6 vertices  
= 2 triangles

start position  
in index buffer

three successive  
vertices in the  
index buffer form  
a triangle

index type

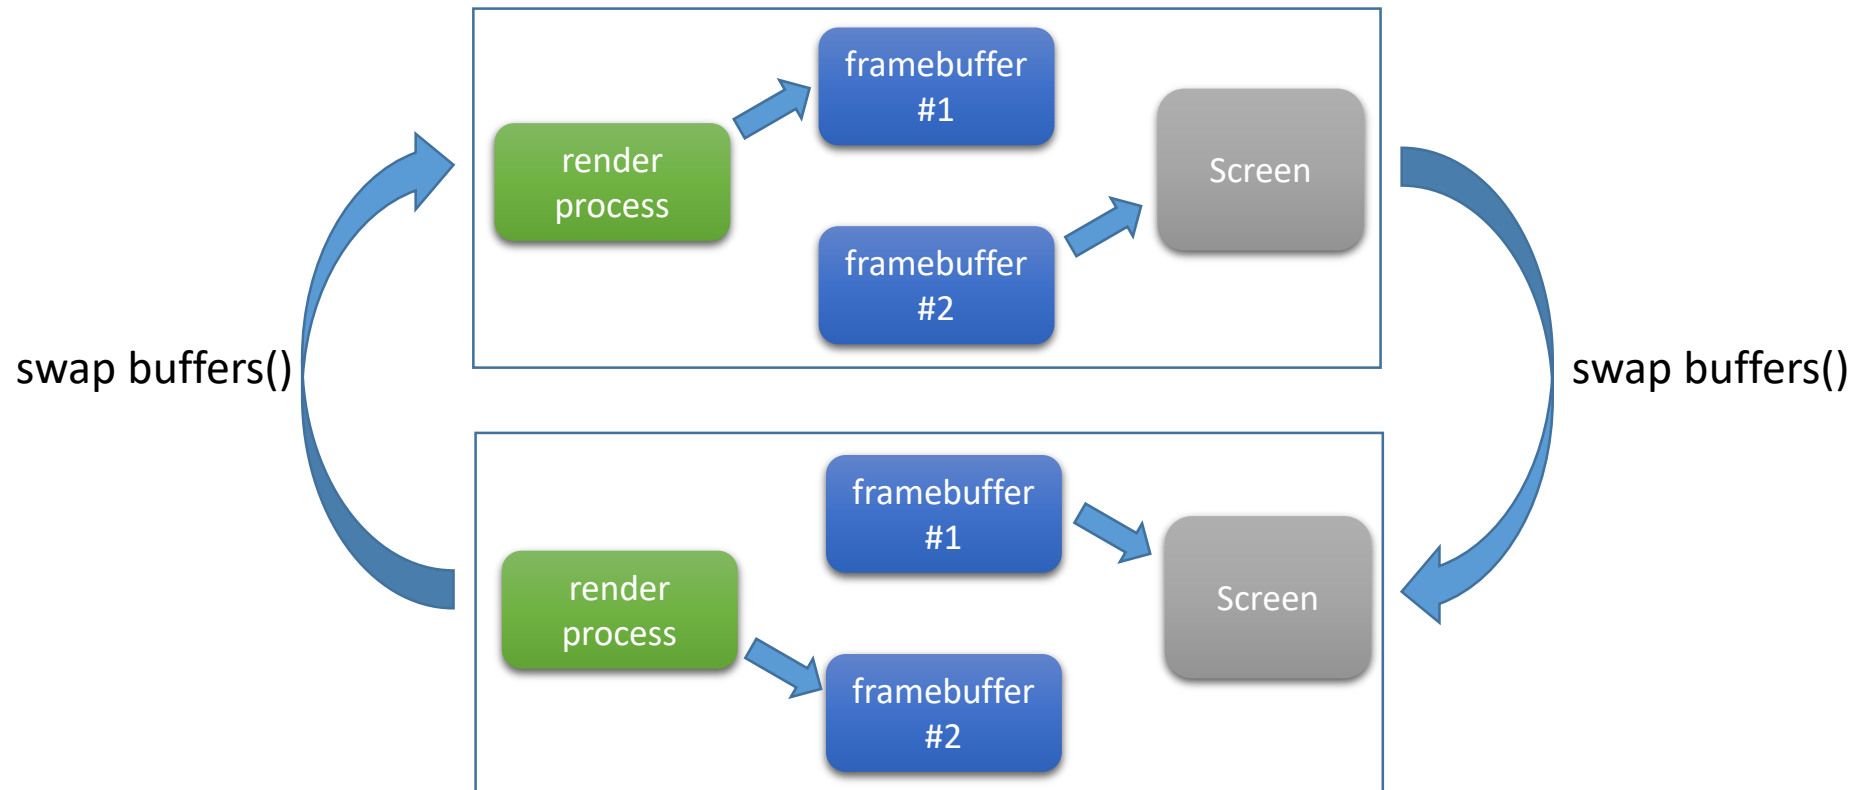
# Let's play

- WebGL



# Render Loop

- To render an animation, we continuously re-render the scene
- Usually two frame buffers available (**double buffering**)
  - One is displayed
  - The other one is being rendered into
  - When new frame has been rendered, buffers are **swapped**



# Render Loop

- Typically the swap is delayed until the next **VSYNC**, i.e. the moment the monitor starts a new frame (usually every  $1/60^{\text{th}}$  second)
- Otherwise: waste of rendering power, images are torn  
→ screen tearing

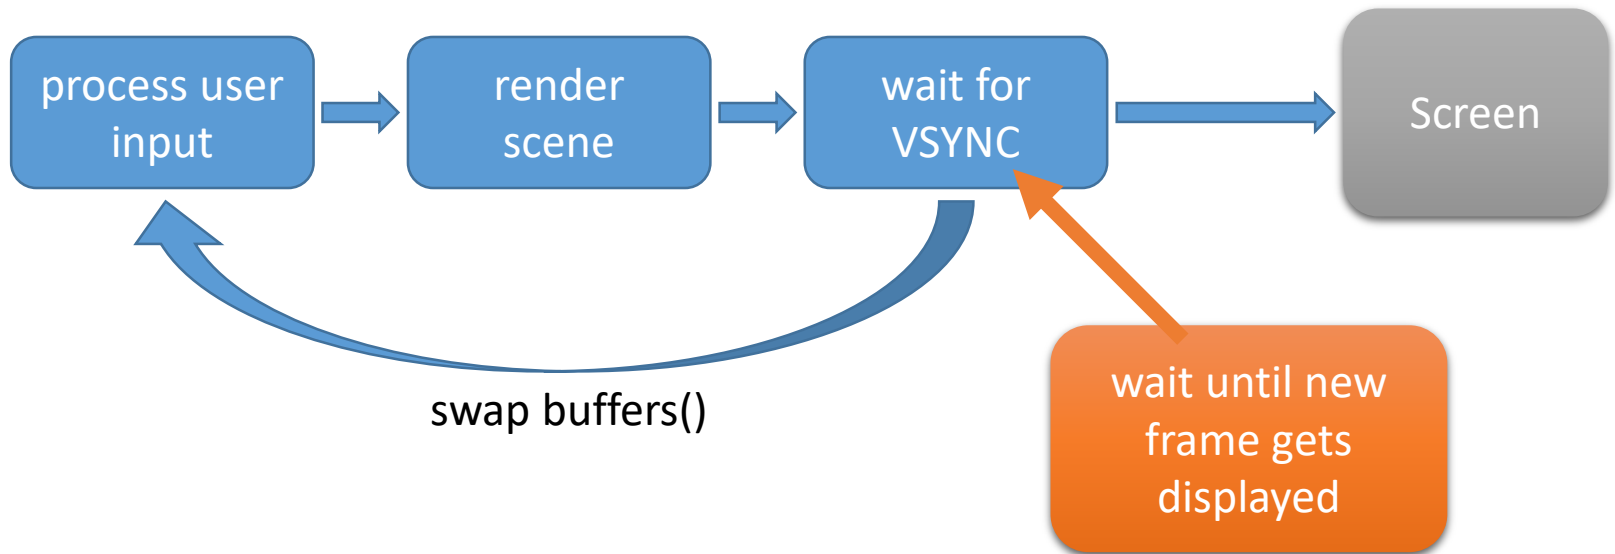


Wikipedia

rendering at 120 Hz,  
display with 60 Hz

# Render Loop

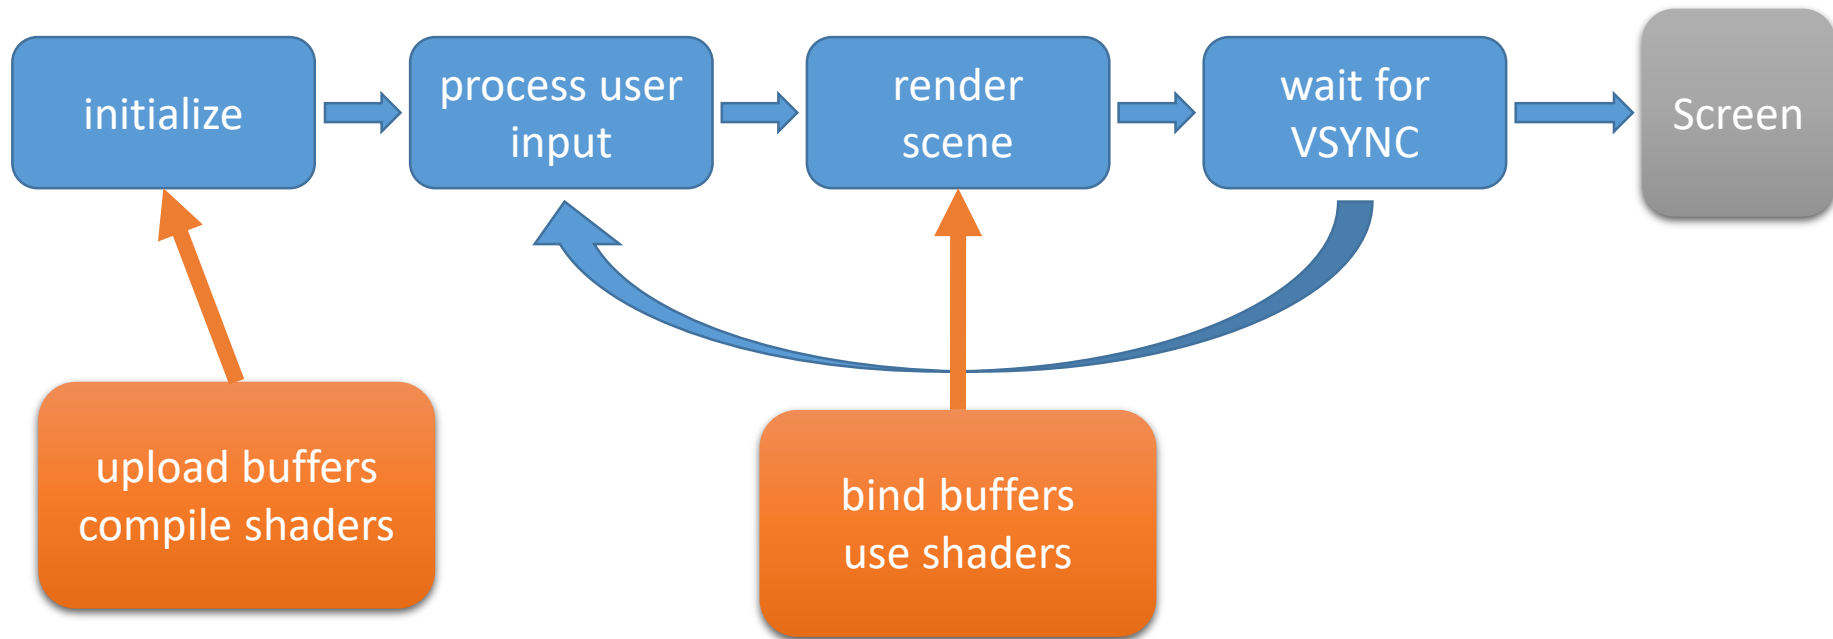
- To render an animation, we continuously re-render the scene
- The screen displays a new image every  $1/60^{\text{th}}$  second
- → **Render Loop**





# Render Loop

- Memory upload, shader compilation and linking etc. are **time expensive**
- **before the render loop**
  - upload buffers (multiple buffers possible)
  - compile and link shaders (multiple shaders / shader programs possible)
- **In the render loop**
  - buffers can be quickly activated by `bindBuffer()`
  - shader programs are activated using `useProgram()`



# Render Loop

- Multiple buffers can be set, as well as multiple shader programs
- `initialize()` (before render loop)

```
var vboObject1 = gl.createBuffer(); // Object 1
gl.bindBuffer(gl.ARRAY_BUFFER, vboObject1);
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(vObject1), gl.STATIC_DRAW);

var vboObject2 = gl.createBuffer(); // Object 2
gl.bindBuffer(gl.ARRAY_BUFFER, vboObject2);
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(vObject2), gl.STATIC_DRAW);
```

- `render()` (within render loop)

```
gl.bindBuffer(gl.ARRAY_BUFFER, vboObject1);
gl.drawElements(gl.TRIANGLES, nVert1, gl.UNSIGNED_SHORT, 0);

gl.bindBuffer(gl.ARRAY_BUFFER, vboObject2);
gl.drawElements(gl.TRIANGLES, nVert2, gl.UNSIGNED_SHORT, 0);
```

# Render Loop

- Same with shaders
- `initialize()`

```
var vertexShader1 = gl.createShader(gl.VERTEX_SHADER);  
var fragmentShader1 = gl.createShader(gl.FRAGMENT_SHADER);  
...  
gl.linkProgram(shaderProgram1);  
...  
var vertexShader2 = gl.createShader(gl.VERTEX_SHADER);  
var fragmentShader2 = gl.createShader(gl.FRAGMENT_SHADER);  
...  
gl.linkProgram(shaderProgram2);
```

- `render()`: object #1 with shader #1, object #2 with shader #2

```
gl.useProgram(shaderProgram1)  
gl.bindBuffer(gl.ARRAY_BUFFER, vboObject1);  
gl.drawElements(gl.TRIANGLES, nVert1, gl.UNSIGNED_SHORT, 0);  
  
gl.useProgram(shaderProgram2);  
gl.bindBuffer(gl.ARRAY_BUFFER, vboObject2);  
gl.drawElements(gl.TRIANGLES, nVert2, gl.UNSIGNED_SHORT, 0);
```

# Next Week

- Transformations
  - Translations, Rotations, Scalings, ...
  - usually happen in the vertex shader