


# Visibility, Culling, Deferred Shading

Computer Graphics  
Winter Term 2016/17


Marc Stamminger / Roberto Grosso

# Occlusion

- Occlusion
  - Essential in 3D graphics



or




- How to create occlusion correctly?
  - Painter's Algorithm
  - Z-Buffer
  - Ray tracing (after Christmas)


# Occlusion

- Painter's Algorithm

- Sort objects from back to front
- Render them in this order
  - front objects draw over back objects
- Very expensive, e.g. sorting of 1 million triangles!
- Cannot handle
  - Penetration



Cyclic occlusion




# Z-Buffer

- Z-coordinates of 3D primitives equal depth values
  - after normalization  $\rightarrow$  z-values are from unit interval  $[-1,1]$
- Interpolate depth value during rasterization, i.e. per pixel depth
  - just as colors for Gouraud-shading
- Z-Buffer
  - buffer with same size as image.
  - Stores depth of currently closest object visible through this pixel
  - Occlusion by simple depth test (z at pixel  $(x, y)$ )  $\rightarrow$  can be implemented in hardware

```
setpixel (x, y, depth, col or)  
    if (zBuffer(x, y) > depth)  
        screen(x, y) := col or  
        zBuffer(x, y) := depth  
    endi f
```


# Z-Buffer

- “Depth” is z-value after normalization  $\Rightarrow z \in [-1,1]$
- Projective mapping maps lines to lines  
→ triangles mapped to triangles by normalization  
(unless triangle intersects  $z = 0 \rightarrow$  problems with clipping after normalization)
- non-perspective interpolation of z-values
- Precision of z-values in z-buffer important
  - depth values mostly close to 1 (comes from perspective mapping)
  - differences in depth become small for distant objects
  - choose n reasonably large
  - at least 24 bit integer or 32 bit float needed



# Z-Buffer

- Tricks: Hidden-Line-Rendering  $\leftrightarrow$  Wireframe Rendering



- render polygons to depth buffer only in 1st pass
- render outlines in 2nd pass and use contents of depth buffer from 1st pass

# Z-Buffer

- Problem: “z-buffer fighting”
  - Pixels from 2nd pass exactly on surface from 1st pass.
  - Effects of rounding
  - Some pixels in 2nd pass occluded
- Solution
  - Move outline towards camera by some delta (or move polygon away from camera)
  - Careful choice of delta required to avoid unwanted additional occlusion effects

# Z-Buffer




wireframe



hidden line with  
polygon offset



Z-fighting problem




polygon and polygon outline  
with polygon offset




# Z-Buffer

- Tricks: Haloing



Gaps for crossing  
lines




- Algorithm

- Render thick outlines to depth buffer only in 1st pass
- Render lines again in normal thickness with offset
- Invisible thick lines hide back lines
- Gaps occur


# Back Face Culling

- For solid objects, every surface triangle has an outer side and an inner side
- Front facing triangle: triangle, of which we see the outer side
- Back facing triangle: non-front facing triangle
- We cannot see back facing triangles (unless we are inside the object)
- But: also front-facing triangles can be occluded (partially or completely)
- → **Back face culling**: remove such backfacing faces




# Back Face Culling

- How can we decide per triangle whether it is back facing ?
- Version 1 (world space)
  - assign a normal  $N$  to each face, pointing outwards
  - render triangle, only iff  $V \circ N > 0$
  - problem: often  $N$  is not known, but only the lighting normal per vertex





# Back Face Culling

- How can we decide per triangle whether it is back facing ?
- Version 2 (screen space)
  - orient vertices
    - when looking from the outside, order vertices counterclockwise
  - when projected to screen space
    - if also counterclockwise in screen space → front face → render
    - if orientation changes → back face → cull



# Back Face Culling

- how do we test orientation?
- Vector product defines orientation
  - given: 3D-vectors  $\mathbf{p}, \mathbf{q}$
  - $\mathbf{r} = \mathbf{p} \times \mathbf{q}$ :
    - $\mathbf{r}$  perpendicular to  $\mathbf{p}$  and  $\mathbf{q}$
    - $\mathbf{p}, \mathbf{q}$  and  $\mathbf{r}$  are “right handed”
- Use this to test orientation of 2D points  $\mathbf{a}, \mathbf{b}, \mathbf{c}$ 
  - lift to 3D:
  - $\mathbf{a} \rightarrow (a_1, a_2, 0)$ ,  $\mathbf{b}, \mathbf{c}$  analog
  - $\mathbf{p} = \mathbf{b} - \mathbf{a}$ ,  $\mathbf{q} = \mathbf{c} - \mathbf{a}$
  - compute  $\mathbf{p} \times \mathbf{q}$
  - $\mathbf{a}, \mathbf{b}, \mathbf{c}$  counterclockwise  
 $\Leftrightarrow (\mathbf{p} \times \mathbf{q})_z > 0$




# Back Face Culling

- supported by OpenGL:
  - // front faces: counter clock wise  
`glFrontFace(GL_CCW);`  
// cull back faces  
`glCullFace(GL_BACK);`  
// back face culling on  
`glEnable(GL_CULL_FACE);`
- Does not replace visibility test!
- It just quickly sorts out 50% of the triangles before rasterization!

# Transparency


- Transparency
  - Technique: Blending
  - During rendering, new pixels do not overwrite previous ones but the values are “blended”




- $\alpha$ -Blending
  - $\text{pixel} := (1 - \alpha) \cdot \text{old} + \alpha \cdot \text{new}$
- Allows drawing of semitransparent objects
  - $\alpha = 0.5 \Rightarrow \text{pixel} := \frac{1}{2} \text{old} + \frac{1}{2} \text{new} \Rightarrow$  half transparent objects
  - $\alpha$  is not transparency but “opacity” =  $1 - \text{transparency}$

# Transparency

- $\alpha$  is often 4th color component  $\Rightarrow$  RGBA instead of RGB
- $(1, 0, 0, 0.1) \Rightarrow$  very transparent red
- $\alpha=1$  corresponds to opaque rendering
  - $(0, 1, 0, 1) \Rightarrow$  opaque green (transparency == 0)
  - $\text{pixel} := 0 \cdot \text{old} + 1 \cdot \text{new} = \text{new} \Rightarrow$  overwrite
- $\alpha$ -blending is not commutative
  - Results change depending on order of blending
  - Rendering without sorting leads to wrong results



50% red over  
50% green over  
100% white




50% green over  
50% red over  
100% white



# Transparency

- Problem: z-Buffer + Transparent objects
  - Example: render 2 semitransparent spheres (1 in front of 2) using  $\alpha$ -blending and z-buffer



- If 1 is drawn before 2, 1 will be opaque because z-buffer hides sphere 2
  - If 2 is drawn first, the result is correct
- Z-Buffer assumes objects are opaque !
- So:
  - Opaque objects should be rendered first with z-buffer
  - Then, transparent objects should be rendered back-to-front


# Transparency

- Simple methods to handle transparency
  - Approach a) (correct)
    - Do not use z-buffer at all but sort objects back to front
  - Approach b) (correct)
    - First render opaque objects with z-buffer
    - Then, “freeze” z-buffer (set to read-only)
    - Finally, sort transparent objects and render back to front
  - Approach c) (faster, but not always correct)
    - First render opaque objects
    - Then, “freeze” z-buffer
    - Finally, render transparent objects without sorting


# Handling Transparency with Depth Peeling

- Idea:
  - Render scene multiple times
  - At each render pass, let only fragments survive that are further away from camera than in previous pass
    - Easy to do in fragment shader
  - Single depth layers of the scene are “peeled”
- Result: Layered Depth Image (LDI) [Shade et al. 98]
  - $n$  depth images, where  $n$  is maximum depth complexity

! not relevant for exam !



# Handling Transparency with Depth Peeling




- Depth peeling from front to back (left to right)
  - Surface: bold black lines
  - Hidden surface: thin black lines
  - “Peeled away” surface: light grey lines

! not relevant for exam !

# Handling Transparency with Depth Peeling

- Transparency can be achieved by blending these layers







Interactive Order-Independent Transparency [Everitt 01]

! not relevant for exam !

# Handling Transparency with Depth Peeling

- Transparency can be achieved by blending these layers



Interactive Order-Independent Transparency [Everitt 01]




! not relevant for exam !

# Deferred Shading

- Pixel shaders often bottleneck
  - Costly shaders
  - Many light sources
  - High depth complexity
    - Draw front-to-back, but sorting within a draw call?
  - Too many fragments (tiny triangles)

# Deferred Shading

- High-depth complexity
  - Occluded fragments will be shaded (overdraw, fill-rate)





UNC Power-Plant  $\approx 50$  mio triangles [Aliaga et al. 99], [Wald et al. 01]



# Deferred Shading

- Idea

- first, fill a **geometry buffer (= “G-Buffer”)**
  - for each pixel, we have all information for the front-most pixel
- Position
- Normal
- Material
- ...



# Deferred Shading

- Then, in a second pass, we compute the lighting for each pixel  
→ only one lighting computation per pixel
- this computation is initialized by rendering a single quad over the entire screen  
→ “Screen space aligned quad”
- or by a **compute shader**


# Render Targets (RT)

- By default OpenGL has
  - Depth Buffer (e.g., 16 / 32 bit)
  - Color Buffer (e.g., RGBA8 / RGBA32F)
- Front and Back Buffer
  - One is bound to pipeline
  - One is set to video out

```
// typical render callback (each frame)
glClear(
    GL_COLOR_BUFFER_BIT |
    GL_DEPTH_BUFFER_BIT
);

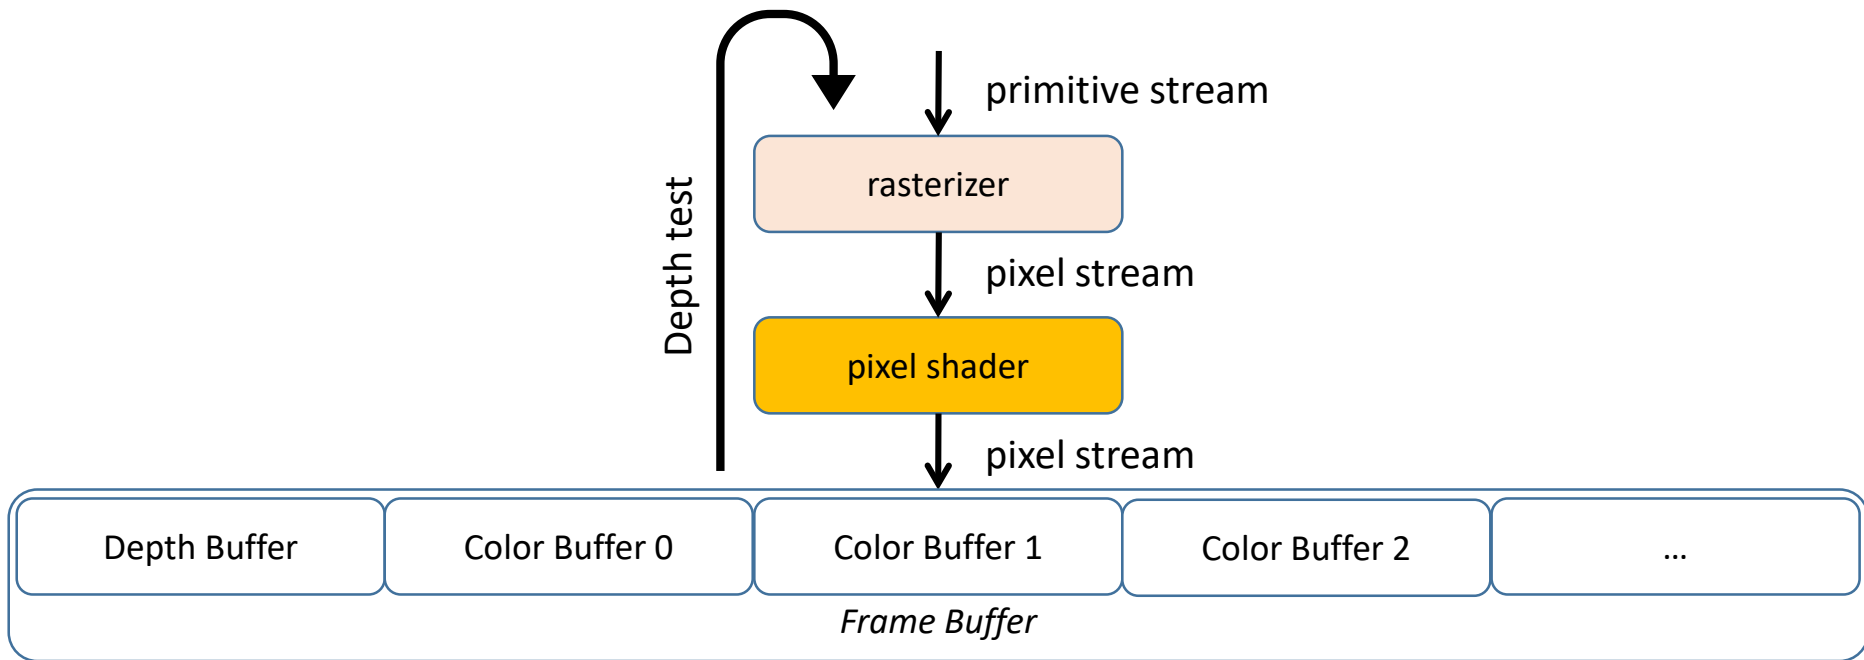
renderScene();

//swap front and back buffer
glutSwapBuffers();
```



```
// fragment shader output
void main(void) {
    ...
    gl_FragColor = color;
    gl_FragDepth = depth;
}
```

# Multiple Render Targets (MRT)



```
// fragment shader output - MRT
void main(void) {
    ...
    gl_FragData[0] = color0;
    gl_FragData[1] = color1;
    gl_FragData[2] = color2;
    ...
    gl_FragDepth = depth;
}
```

# Multiple Render Targets (MRT)

- Up to 8 (color) render targets
  - Can be any format R32F, RGB8, RGBA32F, etc.
  - Must have same dimensions (number of pixels)
  - All RTs must have some number of bits (you can mix channels)
    - OK:
      - RT0 = R32F (32 bit)
      - RT1 = R16G16F (32 bit)
      - RT2 = RGBA8 (32 bit)
    - Does NOT work:
      - RT0 = R16G16F (32 bit)
      - RT1 = RGBA16F (64 bit)
- Only 1 depth buffer
  - There is only one hardware Z-Buffer

# Multiple Render Targets (MRT)

- Render to texture: Framebuffer Objects (FBOs)

```
// FBO: render to texture - create FBO
GLuint color_tex0, color_tex1, depth_tex;
glGenTextures(1, &color_tex0);
...

//allocate texture memory
glBindTexture(GL_TEXTURE_2D, color_tex0);
glTexParameteri(...); //set texture parameters
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGBA8, 256, 256, 0, GL_BGRA, GL_UNSIGNED_BYTE, NULL);
...

//generate framebuffer object
GLuint fbo;
glGenFramebuffers(1, &fbo);
glBindFramebuffer(GL_FRAMEBUFFER, fbo);
glBindFramebufferTexture2D(
    GL_FRAMEBUFFER, GL_DEPTH_ATTACHMENT, GL_TEXTURE2D, depth_tex, 0);
glBindFramebufferTexture2D(
    GL_FRAMEBUFFER, GL_COLOR_ATTACHMENT0, GL_TEXTURE2D, color_tex0, 0);
glBindFramebufferTexture2D(
    GL_FRAMEBUFFER, GL_COLOR_ATTACHMENT1, GL_TEXTURE2D, color_tex1, 0);
...
```

# Multiple Render Targets (MRT)

- Render to texture: Framebuffer Objects (FBOs)

```
// FBO: render to texture - render / bind
glBindFramebuffer(GL_FRAMEBUFFER, fbo);

//will write output into textures
Draw();

glBindFramebuffer(GL_FRAMEBUFFER, 0);
```

- Corresponding fragment shader:

```
// fragment shader output - MRT
void main(void) {
    ...
    gl_FragData[0] = color0;
    gl_FragData[1] = color1;
    gl_FragDepth = depth;
}
```

# G-Buffer

- Typical layout: 16-bit float MRT

RT1	Diffuse.r	Diffuse.g	Diffuse.b	Specular
RT0	Position.x	Position.y	Position.z	Emissive
RT2	Normal.x	Normal.y	Normal.z	Free

- 16-bit float is overkill for diffuse (8-bit would be sufficient)
  - But we need at least 3 x 16-bit for position
  - So don't have a choice due to MRT rules but to make all 16-bit



# Deferred Shading

