

Lecture #15

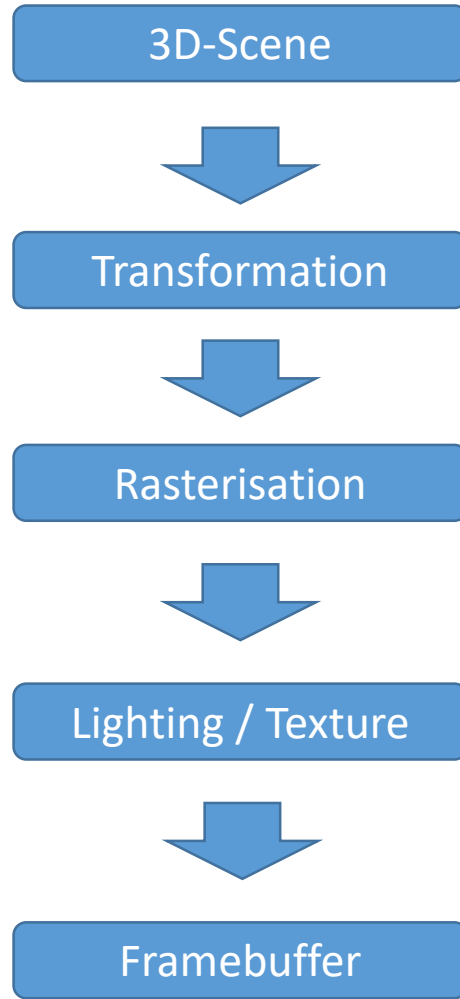
The Rendering Pipeline

Computer Graphics
Winter Term 2016/17

Marc Stamminger / Roberto Grosso

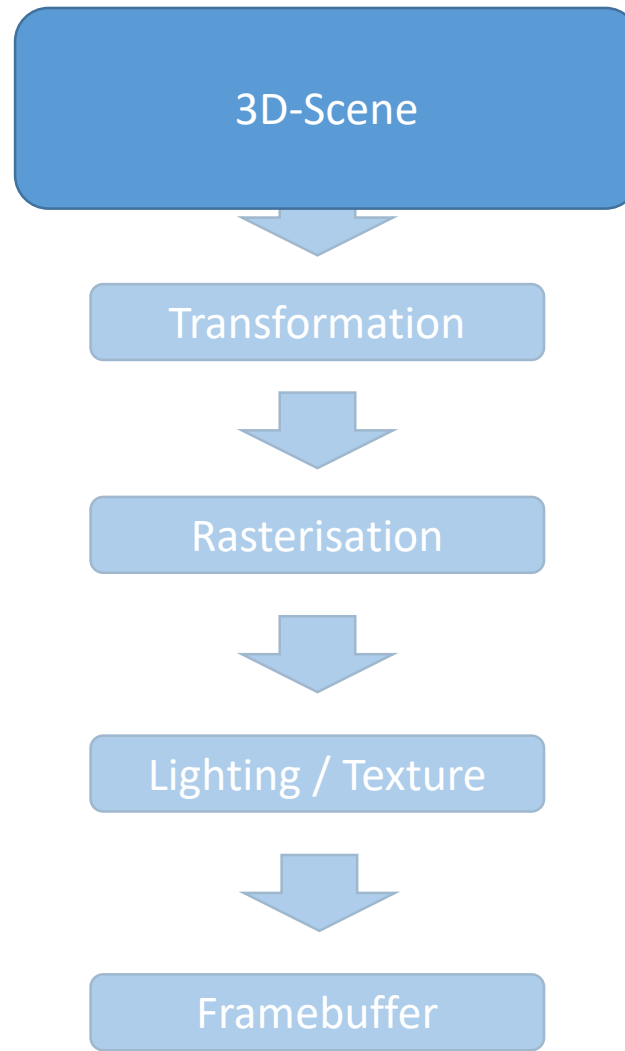
The Rendering Pipeline

- Coarse Version




The Rendering Pipeline

- Coarse Version



3D Objects

- From simple shapes, e.g. a sphere
- Generate using Quad or Triangle Strips
 - enumerate vertices in proper order
 - most vertices appear twice
 - Triangle Fan needed for caps
 - degenerate Strips also possible, where ring at pole collapses to point
 - no index buffer needed




[By MaxDZ8 \(Snapshot from a program I've written.\)](#)
[\[Public domain\], via Wikimedia Commons](#)

```
var v = [...];  
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);  
gl.drawArrays(gl.QUAD_STRIP, 0, v.length);
```

3D Objects

- From simple shapes, e.g. a sphere
- Generate using an Indexed Face Set
 - two buffers needed (vertices and indices)
 - every vertex appears only once




[By MaxDZ8 \(Snapshot from a program I've written.\)](#)
[\[Public domain\], via Wikimedia Commons](#)

```
var v = [...];  
var i = [...];  
  
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);  
  
var ibo = gl.createBuffer();  
gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, ibo);  
gl.bufferData(gl.ELEMENT_ARRAY_BUFFER, new Uint16Array(i), gl.STATIC_DRAW);  
  
gl.drawElements(gl.TRIANGLES, 6, gl.UNSIGNED_SHORT, 0);
```

3D Objects

- From simple shapes, e.g. a sphere
- Generate using an Indexed Face Set **and Strips**
 - two buffers needed (vertices and indices)
 - every vertex appears only once
 - smaller index buffer by using strips
 - strips restart using degenerate triangles:
what does strip ABCDEEFFGHI generate?



[By MaxDZ8 \(Snapshot from a program I've written.\)](#)
[\[Public domain\], via Wikimedia Commons](#)

```
var v = [...];  
var i = [...];  
  
var vbo = gl.createBuffer();  
gl.bindBuffer(gl.ARRAY_BUFFER, vbo);  
gl.bufferData(gl.ARRAY_BUFFER, new Float32Array(v), gl.STATIC_DRAW);  
  
var ibo = gl.createBuffer();  
gl.bindBuffer(gl.ELEMENT_ARRAY_BUFFER, ibo);  
gl.bufferData(gl.ELEMENT_ARRAY_BUFFER, new Uint16Array(i), gl.STATIC_DRAW);  
  
gl.drawElements(gl.TRIANGLE_STRIP, 6, gl.UNSIGNED_SHORT, 0);
```

3D Objects

- From files, e.g. OBJ:

```
v 1.000000 -1.000000 -1.000000
v 1.000000 -1.000000 1.000000
v -1.000000 -1.000000 1.000000
v -1.000000 -1.000000 -1.000000
v 1.000000 1.000000 -1.000000
v 0.999999 1.000000 1.000001
v -1.000000 1.000000 1.000000
v -1.000000 1.000000 -1.000000
vn -0.000000 -1.000000 0.000000
vn 0.000000 1.000000 -0.000000
vn 1.000000 0.000000 0.000000
vn -0.000000 -0.000000 1.000000
vn -1.000000 -0.000000 -0.000000
vn 0.000000 0.000000 -1.000000
f 1//1 2//1 3//1 4//1
f 5//2 8//2 7//2 6//2
f 1//3 5//3 6//3 2//3
f 2//4 6//4 7//4 3//4
f 3//5 7//5 8//5 4//5
f 5//6 1//6 4//6 8//6
```

Vertex positions

Vertex normals

Topology 3//1 means:

vertex with 3rd position
and 1st normal

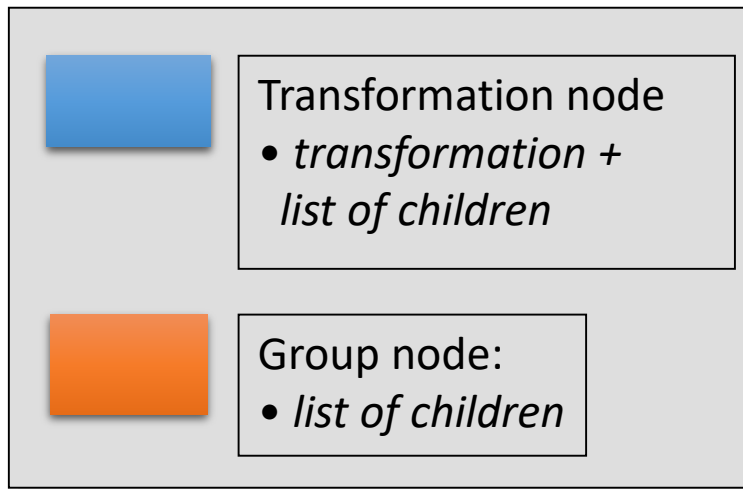
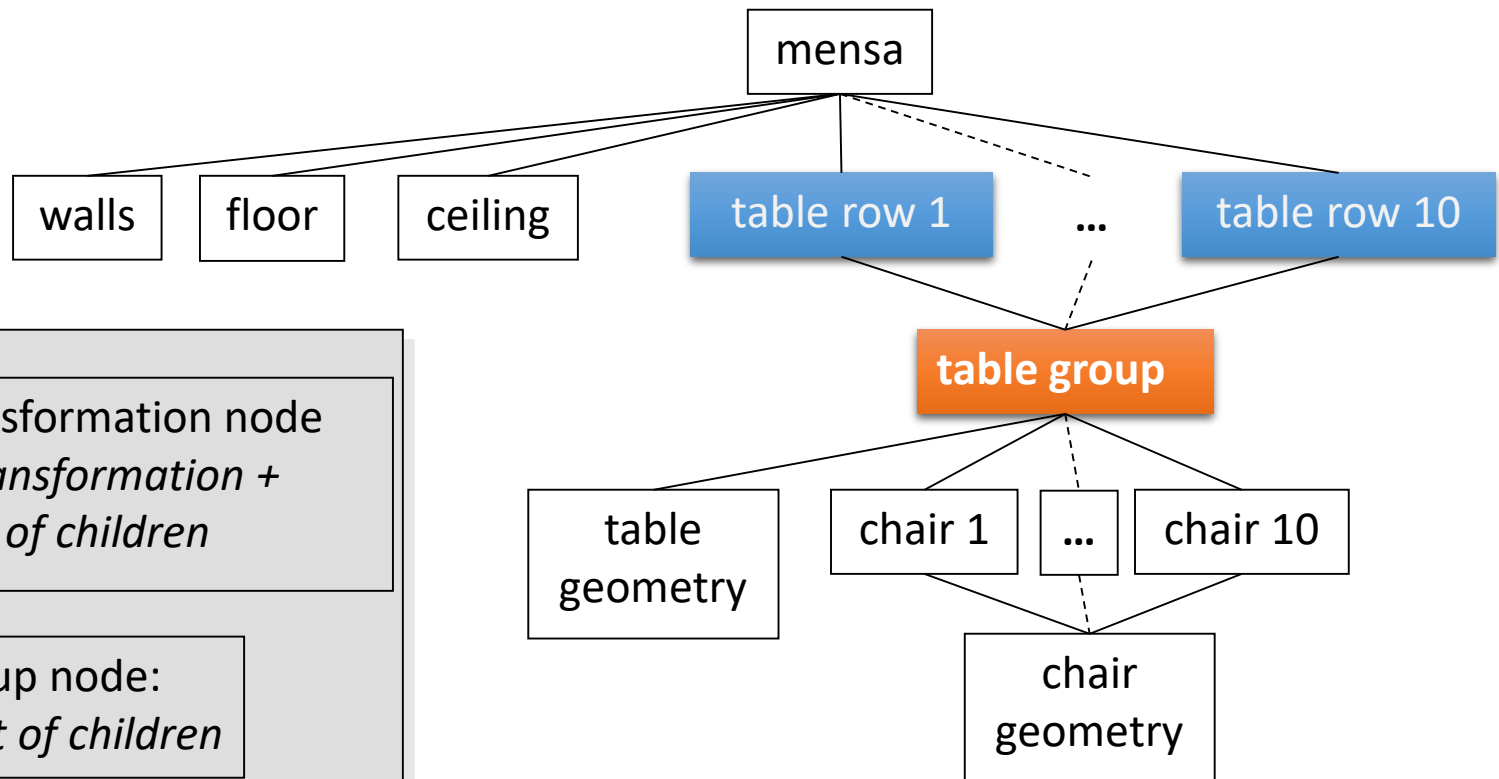
OBJ file format

3D Scenes

- Combine many objects to a 3D Scene
- Each object has
 - material properties
 - a modeling transformation that positions the object in the scene
- **Instancing:**
 - position copies of an object under various transformations
 - supported by OpenGL
- **Scene Graph:**
 - Store the material and transformations in a hierarchy
 - instancing by multiple references to single objects

3D Scenes

- Scene Graph



3D Scenes

- Scene Graph traversal:

- to render a scene graph, we do a depth traversal, always down to the leafes
- on the way, we gather transformations using a matrix stack (or similar)
- at each leaf, we have
 - an object, usually as triangle mesh → bind corresponding buffer
 - its modeling transformation → set as model matrix for shader
 - and material parameters → set shader plus its parameters

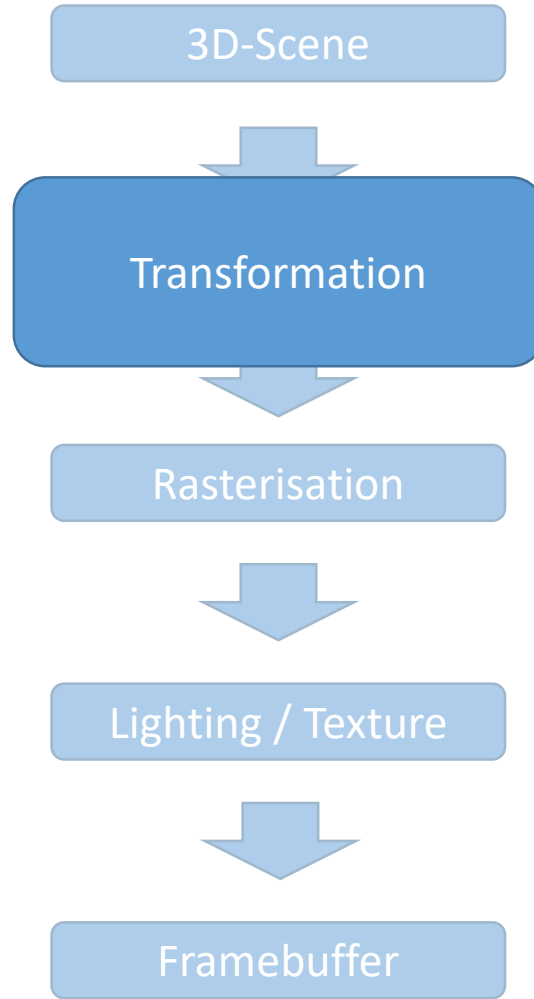
```
class Node { ... }

class TransformNode extends Node {
    Transformation t;
    void render(MatrixStack stack) {
        stack.push(t);
        for (each child i)
            i.render(stack);
        stack.pop();
    }
}

class Object extends Node {
    void render(MatrixStack stack) {
        render Object with modeling matrix stack.top();
    }
}
```

The Rendering Pipeline

- Coarse Version



Transformations


- **Affine Transformations:**

Described using homogeneous coordinates and a matrix

$$\begin{pmatrix} x' \\ y' \\ z' \\ w' \end{pmatrix} = \begin{pmatrix} \vdots & \vdots & \vdots & O_1 \\ e_1 & e_2 & e_3 & O_2 \\ \vdots & \vdots & \vdots & O_3 \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ w \end{pmatrix}$$

- used as

- modeling matrices
- viewing matrix



Transformations

- **Affine Transformations:**

Special case: Rotations

- Representations:

- Orthogonal matrix
- Euler angles (e.g.: yaw, pitch, and roll)

$$R = R_z(\alpha)R_y(\beta)R_z(\gamma)$$


- axis + angle
 - see slides #6
- or quaternions
 - best for interpolation
 - see slides #6

- Modeling Transformation is usually translation + rotation + scale

- Viewing transformation is a translation + rotation

Transformations

- **Viewing:**
 - coordinate transformation to coordinate system aligned with camera
- sets the “extrinsic” camera parameters
- usually defined by
 - camera position (eye)
 - view direction (gaze)
 - up-vector (up)
- or by
 - eye
 - look-at (at)
 - up
 - in this case: gaze = at-eye



Transformations

- Projective Transformations: arbitrary homogeneous matrix

$$\begin{pmatrix} x' \\ y' \\ z' \\ w' \end{pmatrix} = \begin{pmatrix} \ddots & & & \\ & \ddots & & \\ & & \ddots & \\ & & & \ddots \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ w \end{pmatrix} = M \begin{pmatrix} x \\ y \\ z \\ w \end{pmatrix}$$


- Interpretation:
 - Points with $w \neq 0$: points in 3D
 - Points with $w = 0$: points at infinity = directions = vectors
 - \rightarrow first column of M : image of $(1,0,0,0)^T$, second: ...
 - $(1,0,0,0)^T$ is intersection of lines parallel to x-axis
 - if $M_{30} \neq 0$, these parallel lines will intersect in finite space
 \rightarrow parallel lines don't remain parallel
 - columns of M correspond to vanishing points in perspective

Transformations

- very simple perspective = division by z
- as a matrix:


$$M_{perspective} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & -1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

- generates non-linear z : $z \rightarrow 1 - \frac{1}{z}$



Transformations

- Plus selection of window on image plane, plus selection of z-range:




- $M_{perspective}(l, r, b, t, n, f) = \begin{pmatrix} \frac{2n}{r-l} & 0 & \frac{l+r}{r-l} & 0 \\ 0 & \frac{2n}{t-b} & \frac{b+t}{t-b} & 0 \\ 0 & 0 & -\frac{f+n}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{pmatrix}$

Transformations

- Perspective matrix defines “intrinsic” camera parameters:

- field of view (fovy)
 - wide angle lens / tele lens
- aspect ratio (ratio)
 - width over height of image
- near and far plane
 - z-range mapped to $[-1,1]$




- “fovy” defines “top” and “bottom”
- then “aspect” defines “left” and “right”

Transformations

- These three transformations have to be applied in right order:

$$P \ V \ M \begin{pmatrix} x \\ y \\ z \\ w \end{pmatrix}$$

- **M**odeling, **V**iewing, **P**erspective
- Note: right-most matrix is applied **first!**
- M is the first to “see” the point
- Matrix PVM directly transforms a point to the final “canonical view volume”



Transformations

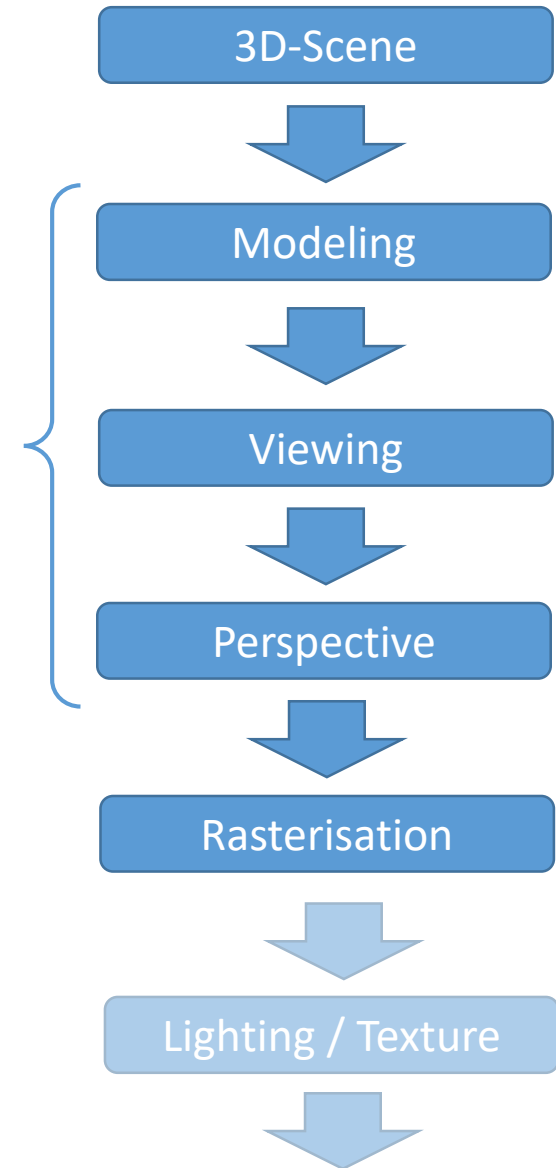
- Transformations are a pipeline on its own:

Transformation

- in OpenGL, these transformations happen in the vertex shader → later
- matrices are passed as **uniforms**

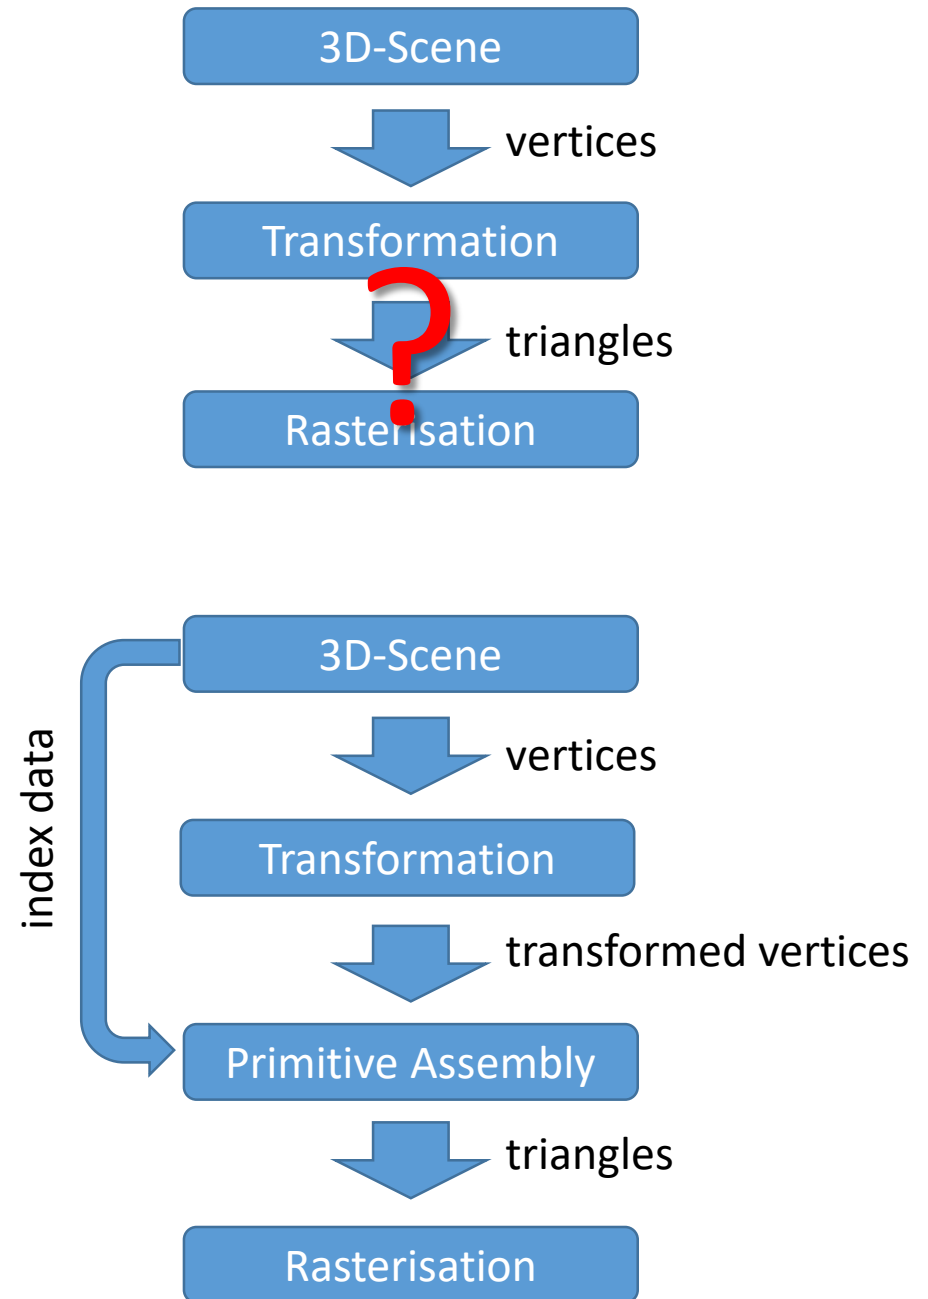
```
// vertex shader - simplest version
attribute vec4 pos;
uniform mat4 PVM;

void main(void) {
    gl_Position = PVM * pos;
}
```



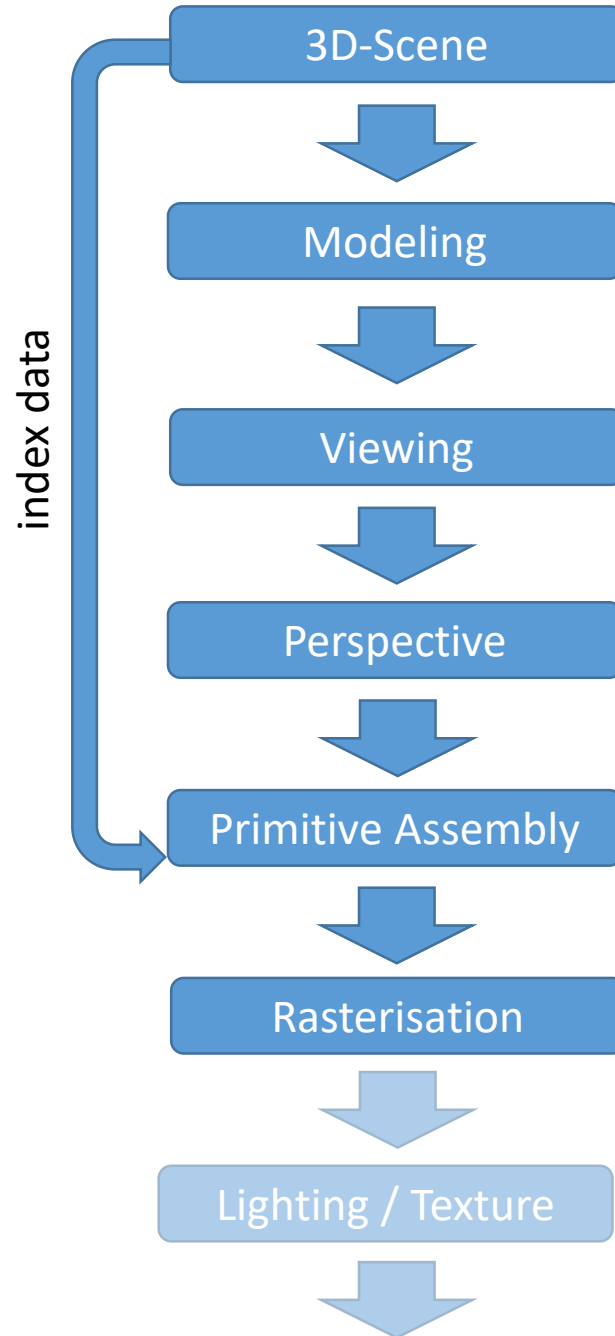
Primitive Assembly

- there is an m:n relation between vertices and triangles
- transformation works on vertices
- rasterization works on triangles
- conversion of vertex stream to triangle stream is done by **primitive assembly**
- This is done either using an implicit topology (triangle strips, fans, ...)
- or using an index buffer



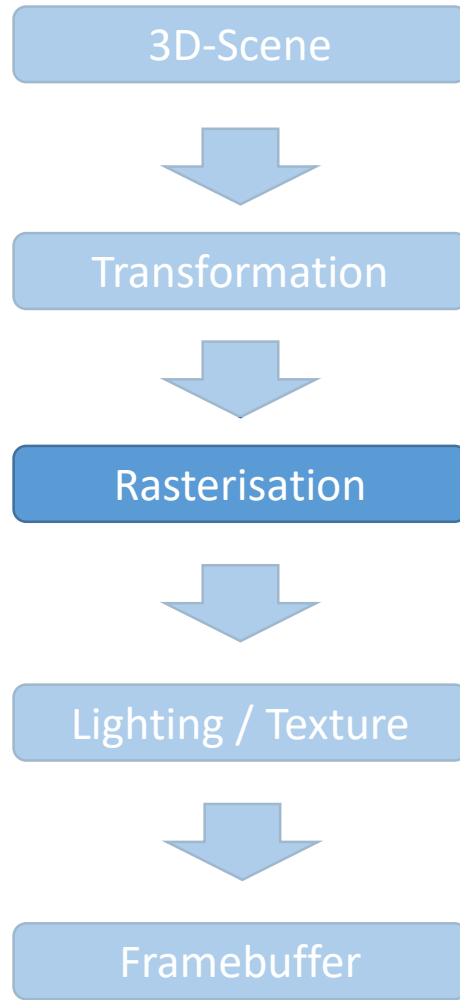
Rendering Pipeline

- so now we have:




Rendering Pipeline

- Coarse Version




Rasterization

- Lines
 - Bresenham




- Polygons
 - Scanline



- Where to get the colors of the set pixels from ?
→ **Lighting**

Lighting

- Given:
 - a number of light sources (point light, parallel light, spot light)
 - a surface point (position and surface normal)
 - and material parameters
 - and the current viewer position
- Compute:
 - color of surface point as seen from the current camera




Lighting

- **Phong Model**


- ambient
- diffuse
- specular




Ambient



Diffuse



Specular




= Phong Reflection

+

+

Shading

- How to integrate lighting into rasterization ?
- All vertices contain a number of attributes
 - at least: vertex position
 - possibly: normal, color, texture coordinates, ...
- all these are interpolated during rasterization
- linear interpolation in screen space
!= linear interpolation in object space
- → **perspective correct interpolation**




Shading

- How to integrate lighting into rasterization ?
- Two important approaches:
- **Gouraud Shading = Vertex Lighting**
Do lighting computation at vertices, then interpolate color
- **Phong Shading = Pixel Lighting**
Interpolate attributes needed for lighting, then compute lighting per pixel
→ better quality, usually more lighting computations → more expensive
- Let's play – Gouraud Shading / Phong Shading / Phong Lighting



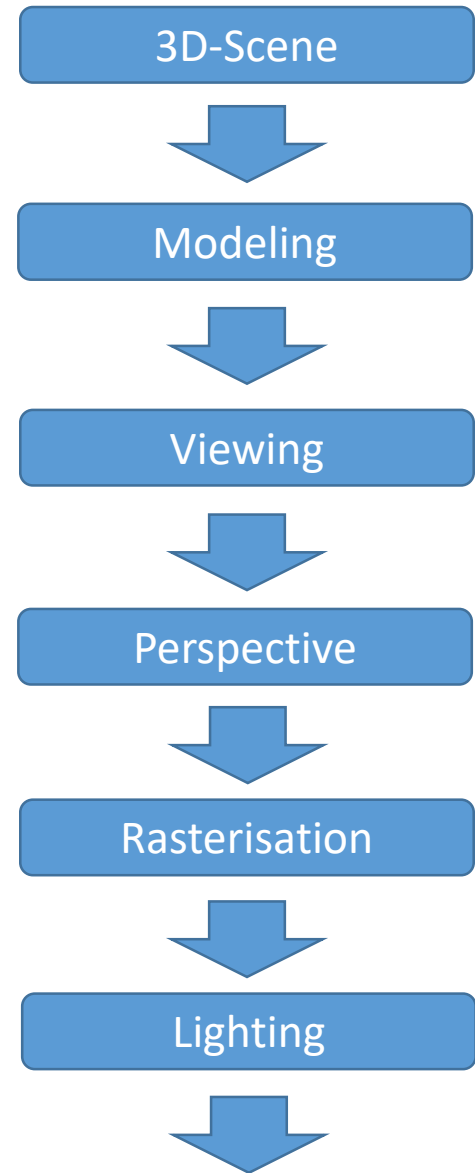
Shading

- Gouraud Shading:
 - Lighting happens per Vertex
 - usually in camera coordinates
 - can also happen in world coordinates (after modeling transformation)






Shading

- Phong Shading
 - Lighting happens after rasterization



Texturing


- Glue image onto objects



- Which part of the image to glue: texture coordinates
- texture coordinates are ordinary attribute that is interpolated during rasterization
→ **perspective correct interpolation!**

Texturing

- For each pixel, the texture value is fetched using the interpolated texture coordinates
 - texture magnification: nearest neighbor or bilinear interpolation (see below)




- texture minification: texture aliasing

minification





magnification



Texturing

- Solution / reduction of texture aliasing: MIPmaps



- Let's play




Depth Buffer


- In addition to attributes, also depth value is interpolated
- Depth buffer used to solve visibility:

```
setpixel (x, y, depth, color)
    if (zBuffer(x, y) > depth)
        screen(x, y) := color
        zBuffer(x, y) := depth
    end if
```

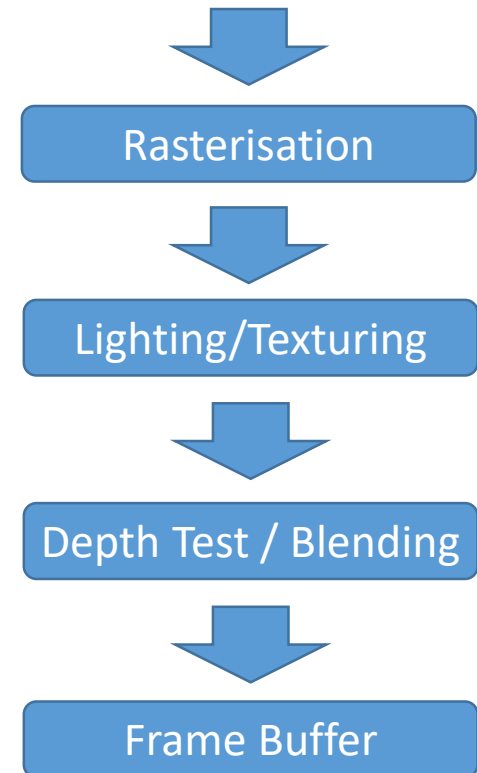
- Blending: combine new pixels with old ones
→ transparency effects



50% red over
50% green over
100% white



50% green over
50% red over
100% white



OpenGL - Shaders

- In OpenGL, some parts of the pipeline are computed using **Shaders**
- Small, C-like programs executed on the GPU
- usually pairs of a vertex and a pixel shader

```
// vertex shader
attribute vec2 pos;
attribute vec3 col;
varying vec3 c;

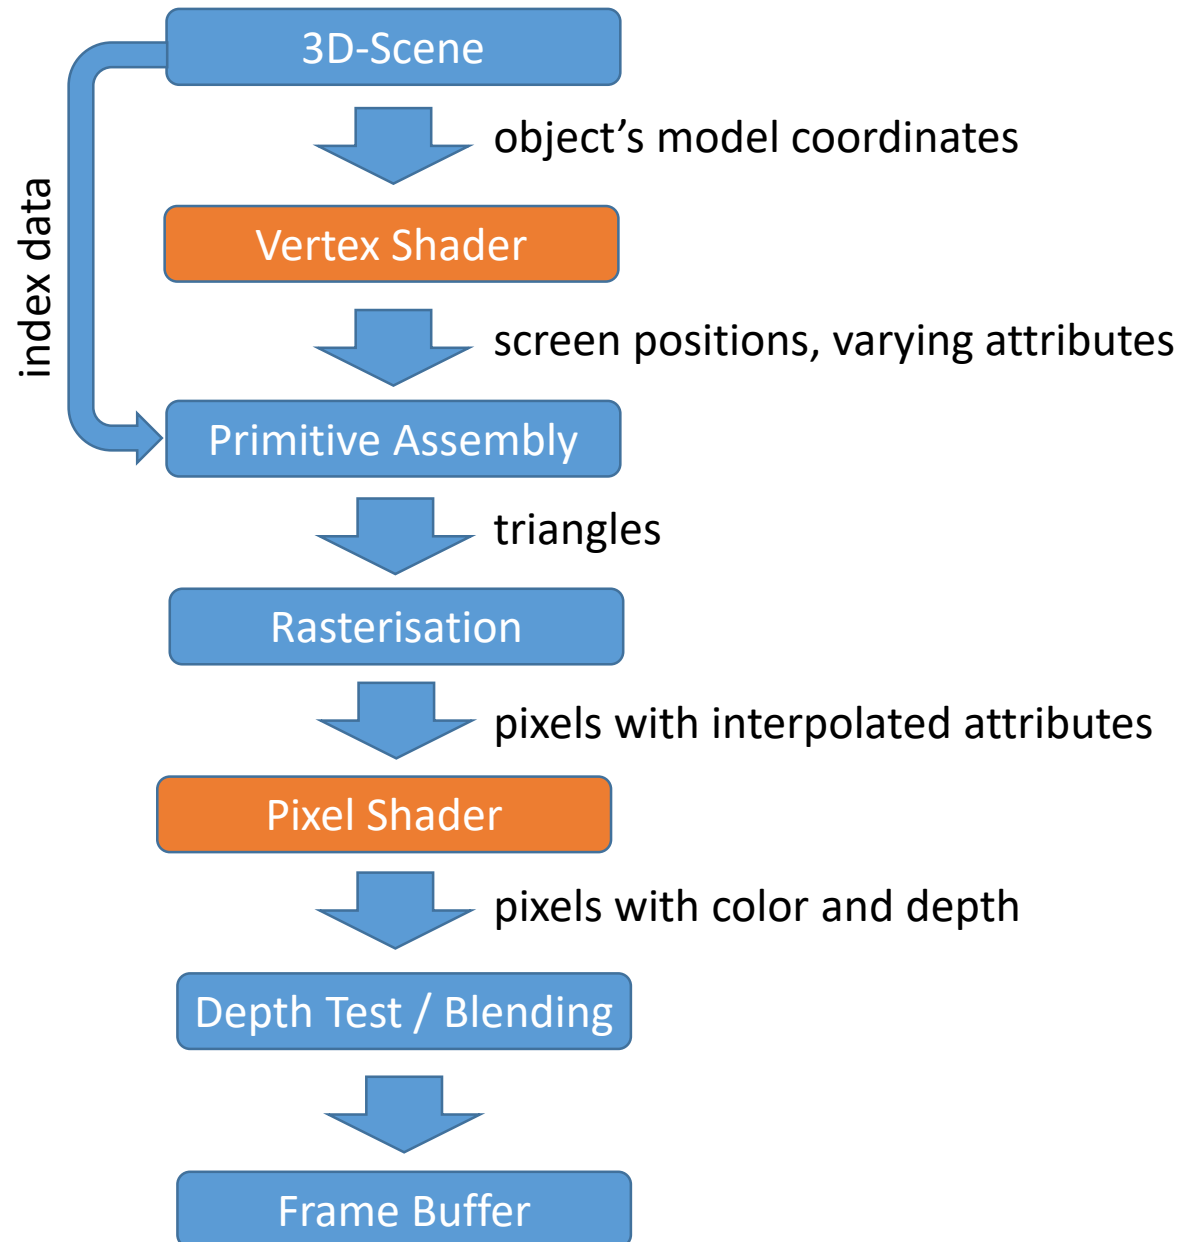
void main(void) {
    gl_Position = vec4((pos+offset)*zoom, 0.0, 1.0);
    c = col;
}
```

```
// fragment shader
precision highp float;
varying vec3 c;

void main(void) {
    gl_FragColor = vec4(c,1);
}
```

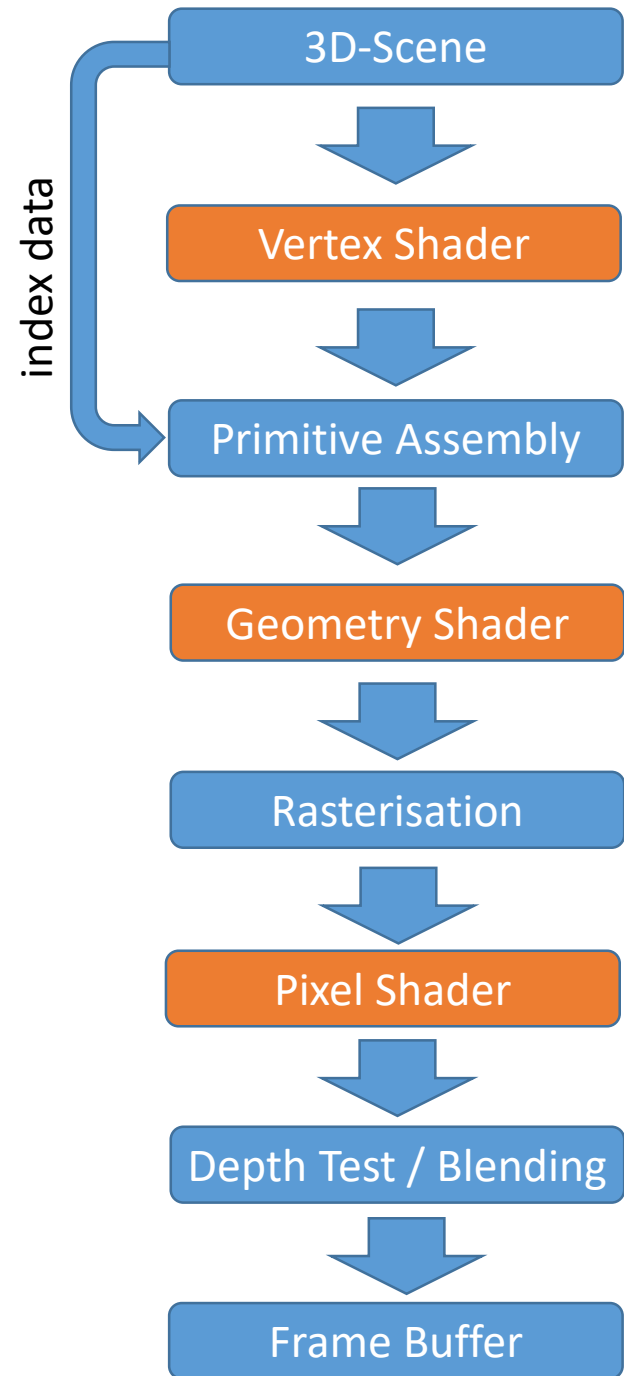
OpenGL - Shaders

- OpenGL pipeline



OpenGL - Shaders

- with geometry shader
(handled in advanced exercises)



Merry Christmas !