

Lecture #18

# Ray Tracing – Acceleration Structures

Computer Graphics  
Winter Term 2016/17

Marc Stamminger / Roberto Grosso

# Introduction

- Billions of rays are required to generate high quality images
- Two major components have to be optimized
  - Basic primitive tests, e.g. ray-triangle intersection  
→ last lecture
  - strategies and data structures to reduce the number of necessary tests  
→ this lecture

# Ray tracing Complexity

- Plain Ray Casting:
  - for each of the  $n_p$  pixels
    - test eye ray against each of the  $n$  scene primitives
  - $O(n_p n)$  is enormous!
  - typically  $n_p, n > 1.000.000$
  - 1 mio objects, 1 mio pixels  $\rightarrow$  1 trillion ( $10^{12}$ ) intersection tests...
- Ray Casting: 1.000 secondary rays per pixel common
  - $\rightarrow 10^{15}$  intersection tests
- But: with acceleration structures the inner loop can be accelerated to  $O(\log n)$ 
  - $\rightarrow$  entire algorithm gets  $O(n_p \log n) \rightarrow$  practical!

# Ray tracing – Acceleration Techniques

- Fact: 90% of the time goes into intersection tests
- Strategies for speeding up ray tracing:
  - Bounding Volumes
  - Space Partitioning
  - Ray coherence (trace bundles of rays)
- Requires one preprocessing step (for all rays)


# Acceleration Techniques

- Bounding volumes (BV):
- Find (geometrically simple) surrounding volumes for the complex objects
  - Spheres
  - bounding boxes
- Choose BV such that the intersection test is simple and efficient.

# Acceleration Techniques

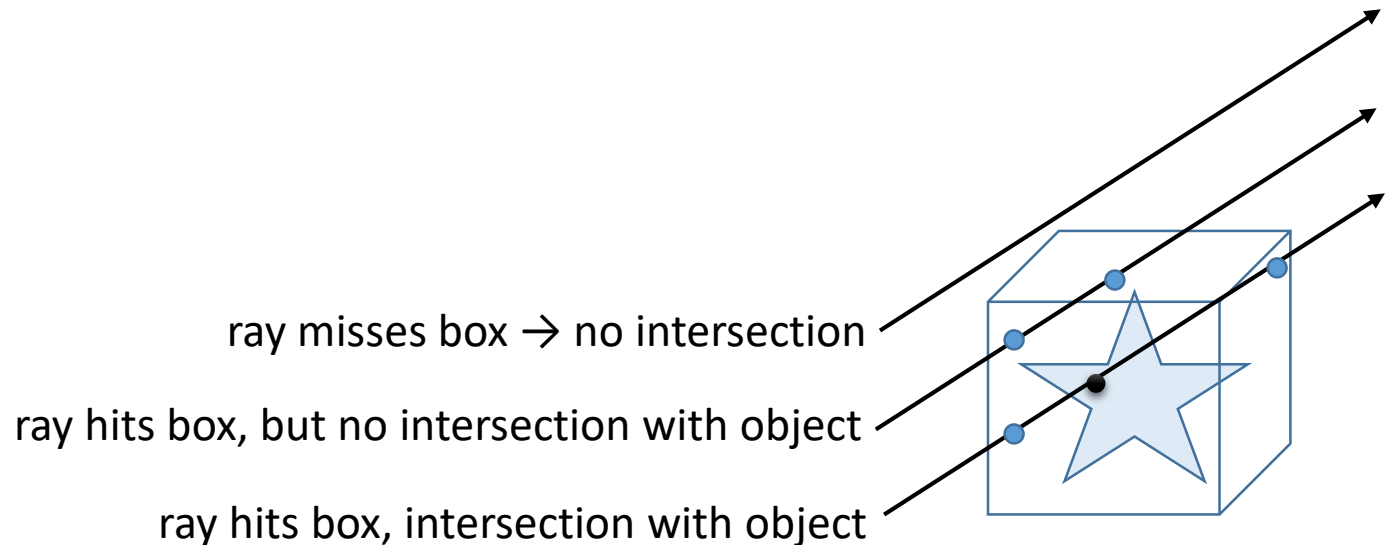
- Bounding volumes (BV) – Intersection test

```
if (intersect (ray, BV) == true) then  
    intersect(ray, BV.objects);  
end if
```




# Ray – Box Intersection

- Another important test is the intersection of a ray with a box
  - box can be a scene primitive
  - but mostly important for accelerated ray – scene interaction tests
- Idea:
  - compute a box surrounding a complex object
  - if ray misses this bounding box, no tests with complex object necessary



# Ray – AABB Intersection


- AABB: axis aligned bounding box
  - box is aligned with the main axes
  - Intersection of three **slabs**  $[x_{min}, x_{max}]$ ,  $[y_{min}, y_{max}]$ ,  $[z_{min}, z_{max}]$





# Ray – AABB Intersection

- Intersection test
  - In 3D a point is inside the AABB if and only if it is inside all the three slabs
  - a ray intersects the AABB if and only if the intersection segments of the ray with the three slabs are overlapping




# Ray – AABB Intersection


- Intersection test with ray


$$p(t) = e + td$$
$$t_n = \frac{d_n - e \circ n}{d \circ n}$$
$$t_f = \frac{d_f - e \circ n}{d \circ n}$$

- Intersection with  $s_1$  is  $[t_1, t_2]$   
Intersection with  $s_2$  is  $[t_3, t_4]$   
→ ray intersects iff  $[t_1, t_2] \cap [t_3, t_4] \neq \{\}$



# Ray – AABB Intersection


$$t \in [t_{xmin}, t_{xmax}]$$


$$t \in [t_{ymin}, t_{ymax}]$$


$$t \in [t_{xmin}, t_{xmax}] \cap [t_{ymin}, t_{ymax}]$$


$$t_{entry} = \max\{t_{xmin}, t_{ymin}\}$$

$$t_{exit} = \min\{t_{xmax}, t_{ymax}\}$$



# Ray – AABB Intersection

- Representation of an AABB

```
// region R = { (x, y, z) | min.x<=x<=max.x,  
// min.y<=y<=max.y, min.z<=z<=max.z }  
struct AABB {  
    Point min;  
    Point max;  
};
```

# Ray – AABB Intersection


```
// Intersect ray  $R(t) = p + t \cdot d$  against AABB a. When intersecting,  
// return intersection distance tmin and point q of intersection  
int IntersectRayAABB(Point p, Vector d, AABB a, float &tmin, Point &q)  
{  
    tmin = 0.0f;           // set to -FLT_MAX to get first hit on line  
    float tmax = FLT_MAX; // set to max distance ray can travel (for segment)  
  
    // For all three slabs  
    for (int i = 0; i < 3; i++) {  
        if (Abs(d[i]) < EPSILON) {  
            // Ray is parallel to slab. No hit if origin not within slab  
            if (p[i] < a.min[i] || p[i] > a.max[i]) return 0;  
        } else {  
            // Compute intersection t value of ray with near and far plane of slab  
            float ood = 1.0f / d[i];  
            float t1 = (a.min[i] - p[i]) * ood;  
            float t2 = (a.max[i] - p[i]) * ood;  
            // Make t1 be intersection with near plane, t2 with far plane  
            if (t1 > t2) Swap(t1, t2);  
            // Compute the intersection of slab intersections intervals  
            tmin = Max(tmin, t1);  
            tmax = Min(tmax, t2);  
            // Exit with no collision as soon as slab intersection becomes empty  
            if (tmin > tmax) return 0;  
        }  
    }  
  
    // Ray intersects all 3 slabs. Return point (q) and intersection t value (tmin)  
    q = p + d * tmin;  
    return 1;  
}
```

# Ray – AABB Intersection

- The test is a special case of the intersection test of ray against a Kay-Kajiya slab volume, T. Kay and J. Kajiya, SIGGRAPH 1986
- The Kay-Kajiya test is a specialization of the Cyrus-Beck clipping algorithms, M. Cyrus and J. Beck, Computer and Graphics, 1978


# Acceleration Techniques

- Bounding Boxes



# Acceleration Techniques



- Bounding Box Hierarchies





# Acceleration Techniques

- BVH: Bounding Volume Hierarchy



# Acceleration Techniques

- BVH (Bounding Volume Hierarchy)
  - adapts well to arbitrary geometries
  - memory consumption is predictable
  - each geometric primitive (e.g. triangle) occurs in exactly one leaf node
  - nodes can overlap in space
  - recursive traversal algorithm (use a stack!)
  - according to splitting strategy, objects can be split when straddling the subdivision plane

# Bounding Volume Hierarchy

- Construction
  - Recursive construction algorithm
  - Initialize root node. Fill it with all triangles in the scene
  - Recursively subdivide the root node
  - Stop recursion if depth of node exceeds a given value or the node contains less than a given number of triangles
  - Split the node into a left and right child otherwise

# Bounding Volume Hierarchy

- Construction

```
buildTree(Scene& scene) {  
    create root node containing all triangles;  
    subdivide(root);  
}  
  
subdivide(Node& node) {  
    if (node.depth == MAX_DEPTH ||  
        node.numTriangles <= MIN_TRIANGLES)  
        return;  
    compute optimal split position;  
    create left and right child nodes;  
    sort triangles into left and right nodes;  
    subdivide(left);  
    subdivide(right);  
}
```

# Bounding Volume Hierarchy


- Computing the optimal splitting plane
  - spatial median: split nodes in the middle along axis with largest extent
  - object median: split nodes so that left and right children contain the same number of triangles
  - cost function: minimize a cost function

# Bounding Volume Hierarchy

- Surface Area Heuristics, SAH

Theorem:

- given a box  $B$  completely contained in a box  $A$
- the probability that a ray traversing  $A$  intersects  $B$  is given by  $SA(B)/SA(A)$ ,  $SA(\cdot)$  is the surface area.



# Bounding Volume Hierarchy

- Cost function for a split

$$C = C_T + |P_l| \frac{SA(B_l)}{SA(B_p)} + |P_r| \frac{SA(B_r)}{SA(B_p)}$$

- where  $P_l$  and  $P_r$  are the number of primitives in the left and right child nodes respectively, and  $B_l$ ,  $B_r$  and  $B_p$  are the left, right and parent bounding boxes.  $C_T$  is the cost of computing a ray-primitive intersection relative to the cost of traversing a node.


# Bounding Volume Hierarchy

- Problem statement
  - find the split position which minimizes the cost function
- Computing the SAH
  - Difficult task, the search space is extremely large!
  - [Ingo Wald: “On fast Construction of SAH-based Bounding Volume Hierarchies”, 2007](#)
- Computing the SAH - Fast approximate SAH construction
  - generate candidate split locations for X,Y and Z
    - use minima and maxima bounding boxes of all geometric primitives
  - Evaluate SAH cost function at each location
  - select splitting plane with lowest cost
  - Compute  $P_l$  and  $P_r$  incrementally




# Acceleration Techniques

## Uniform Space partitioning



# Acceleration Techniques

- Uniform Grid



# Acceleration Techniques

- Uniform Grid


- simple and fast traversal algorithm (e.g. line rasterization in 3D, see later on)
- Not appropriate to handle geometry which is not equally distributed in space, high cost stepping empty cells, cannot skip empty space
- Inadequate for handling geometries of very different sizes, no optimal cell size

# Ray traversal: uniform grid

- Intersection test with a uniform grid
  - two strategies
    - Extended 3D-Bresenham line drawing (rasterization)
    - follow ray from cell boundary to cell boundary



# Ray traversal: uniform grid

- Extended Bresenham
  - major problem: rasterization method will miss some cells
  - Can be corrected by carefully looking at decider variable




# Ray traversal: uniform grid

- Alternative method (Amanatides 1987)
  - step from cell boundary to cell boundary
  - Key idea:
    - distance between vertical boundaries is constant, the same applies for horizontal boundaries.
    - if the ray crosses a vertical boundary, step along the x axis; if the ray crosses an horizontal boundary, step along the y axis.



# Ray traversal: uniform grid



# Ray traversal: uniform grid

- Method
  - maintain two variables measuring distance to next vertical and horizontal planes
  - if distance  $t_x$  to vertical plane is less than the distance  $t_y$  to horizontal plane step along x axis, else step along y axis.
  - Updates:  $t_x += \Delta t_x$  and  $t_y += \Delta t_y$
- Given a ray  $eye + t \cdot dir$  with  $t_{min}$  and  $t_{max}$ 
  - If ray has direction  $(d_x, d_y, d_z)$  then  $\Delta t_x = \frac{M}{d_x}, \Delta t_y = \frac{M}{d_y}, \Delta t_z = \frac{M}{d_z}$
  - where M is the grid size
- Initialization:
  - Find grid cell of ray's starting point  $eye + t_{min} * dir \rightarrow (i_x, i_y, i_z)$
  - Find ray parameters  $t_x, t_y, t_z$  where ray leaves corresponding slabs



# Ray traversal: uniform grid

```
float M = grid_cell_size;

traverseGrid(float3 eye, float3 dir, float tmin, float tmax)
{
    // first grid cell (ix,iy,iz) depending on eye + tmin*dir
    int ix = ..., iy = ..., iz = ...;

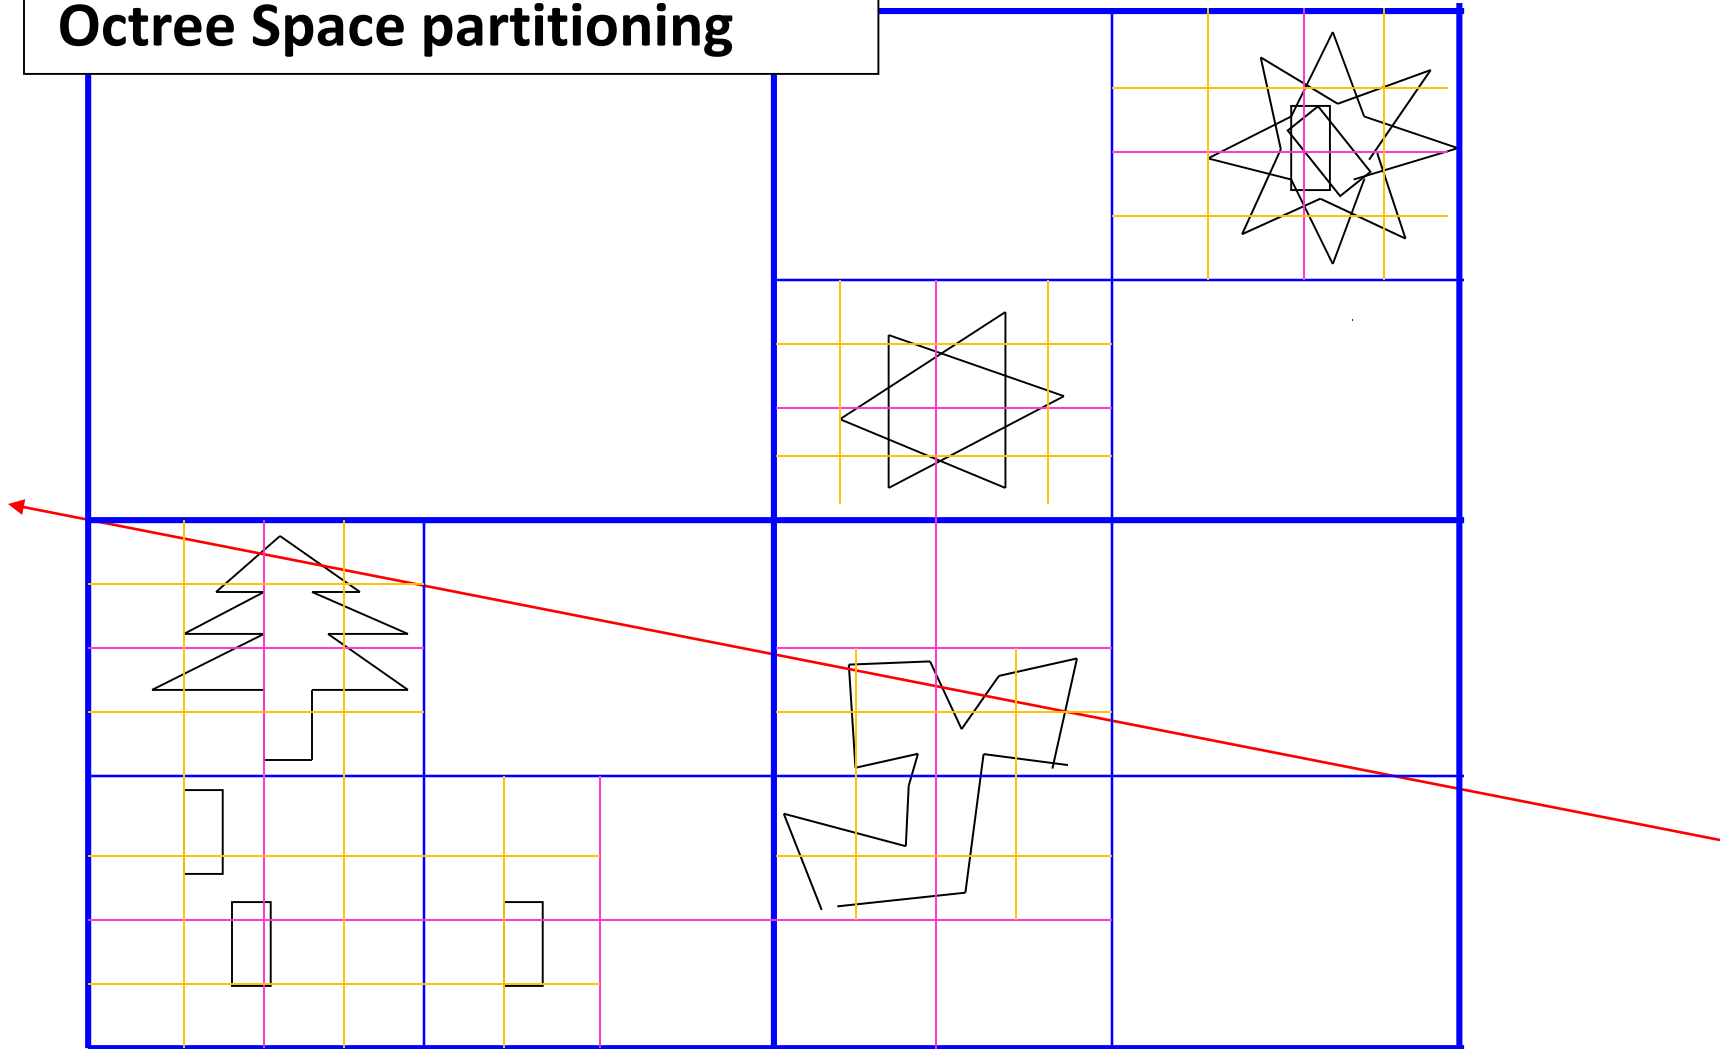
    // ray parameter along x,y,z leaving the corresponding slab
    float tx = ..., ty = ..., tz = ...;

    // step size for ray parameter along x,y,z
    float dtx = M / dir.x, dty = M / dir.y, dtz = M / dir.z;

    while (tx < tmax || ty < tmax || tz < tmax) {
        if (tx < ty && tx < tz) { // go along x
            ix++;
            intersectWithCell(eye, dir, tmin, tmax, ix, iy, iz);
            tx += dtx;
        } else if ... // other dimensions analog
    }
}
```

# Acceleration Techniques



## Octree Space partitioning



# Acceleration Techniques

- Octree Space Partitioning
  - Cannot adapt perfectly to scene (split planes fixed)
  - traversal of children not very efficient
- Better alternative: kd-trees

# Kd-Tree



# Kd-Tree


- Properties
  - Space subdivision, i.e. leaf nodes do not overlap
  - Every node is subdivided spatially in one dimension only (x, y, and z)
  - Split plane can be chosen freely → adapts well to arbitrary geometr
  - Split dimension changes over levels
  - Geometric primitives (e.g. triangles) may occur in more than one leaf node  
→ contrary to BVHs !
  - nodes do not overlap  
→ contrary to BVHs !
  - very efficient recursive traversal

# Kd-Tree

- Leaf-storing tree
  - internal nodes store dividing plane (splitting axis and splitting position) and reference to children nodes
  - leaf nodes stores triangles intersecting the half spaces



# Kd-Tree

- Example
  - assign triangle to both left and right children nodes



# Kd-Tree

- Example







# Kd-Tree

- Traversal: simple approach
  - trace ray by traversing the kd-tree and intersecting ray with splitting planes
  - line segment (ray)  $p(t) = e + t d; 0 \leq t < t_{max}$

# Kd-Tree

- Method

- intersect line segment against node's splitting plane  $\rightarrow t$
- test if solution satisfies  $t \in [0, t_{max})$ 
  - if  $t \in [0, t_{max})$  descend both children recursively
    - descend children in ray order (depending on sign of direction)
      - $\rightarrow$  find closest hit points first
      - $\rightarrow$  allows for early exit
  - if  $t \notin [0, t_{max})$  descend child containing  $t = 0$



# Kd-Tree

- Traversal
  - The ray is traced by traversing the kd-Tree. Three stacks are used for this purpose
    - nodeStack for kd-tree nodes
    - tMinStack for parameter values at the entry points
    - tMaxStack for parameter values at the exit points

# Kd-Tree

- Traversal, cont'd
  - Stacks are initialized at the root node, intersecting the bounding box of the complete scene
  - Following two cases has to be considered
    - internal nodes
    - leaf nodes
  - Stop traversal if intersection was found or stacks are empty

# Kd-Tree

```
// test if ray intersects scene bounding box
bBoxIsect = intersect(ray, scene.boundingBox);
if (!bBoxIsect.hit)
    return NO_INTERSECTION;
nodeStack.push(kdTree.rootNode);
tMinStack.push(bBoxIsect.tMin);
tMaxStack.push(bBoxIsect.tMax);

// main loop
while (!stack.empty()) {
    node = nodeStack.top(); nodeStack.pop();
    tMin = tMinStack.top(); tMinStack.pop();
    tMax = tMaxStack.top(); tMaxStack.pop();


    // handle internal and leaf cases
    if (node.type == LEAF_NODE)
        handle leaf node;
    else
        handle internal node;
}
```

# Kd-Tree

- tMin and tMax
  - tMin is the parameter value along the ray at the entry point into the current node
  - tMax is the parameter value at the exit point
  - tMin and tMax are initialized by intersecting the ray with the scene bounding box
  - If no intersection was found, the traversal returns immediately


# Kd-Tree

- tMin and tMax



# Kd-Tree

- Leaf Nodes
  - intersect ray with geometric primitives in node.
  - if an intersection was found, return the closest intersection point and stop traversal
  - continue traversal otherwise







# Kd-Tree

- Internal Nodes
  - determine near and far children
  - cases
    - ray intersect near child only: push near child onto stack
    - ray intersect far child only: push far child onto stack
    - ray intersect both children: push far child, then near child (near child will be processed first)
  - continue traversal


# Kd-Tree



**$\text{ray.direction.x} > 0$**   
Near child: left  
Far child: right



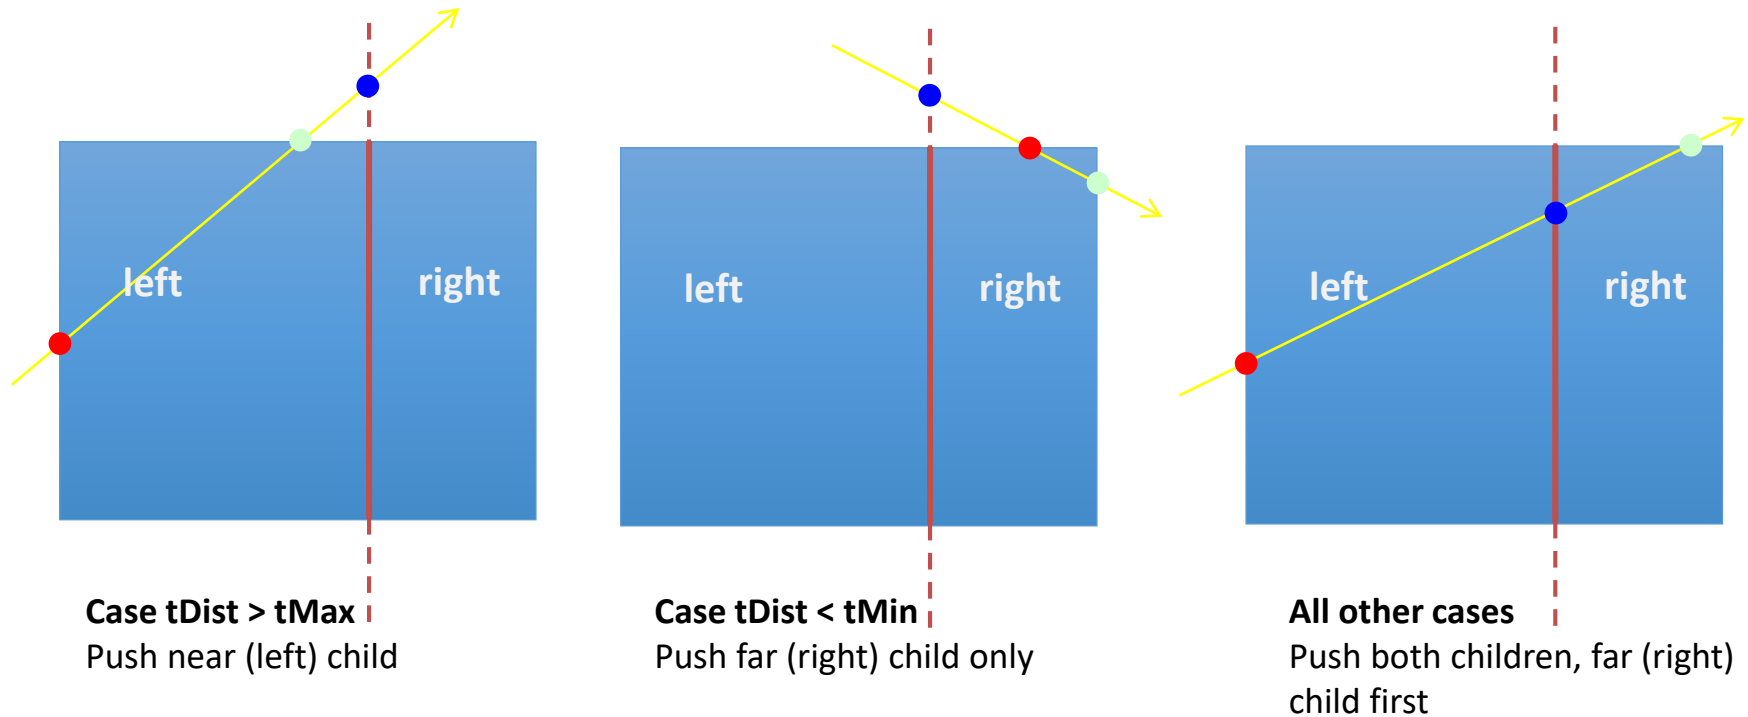
**$\text{ray.direction.x} < 0$**   
Near child: right  
Far child: left



**$\text{ray.direction.x} == 0$**   
Special case during traversal

# Kd-Tree



- Internal Nodes: pushing the stack



- tMin, entry point
- tMax, exit point
- tDist, intersection with Kd plane

# Kd-Tree

- Internal Node: special cases
  - ray is parallel to splitting plane
  - near child depends on position of ray starting point
- Example:  
split axis is x



# Acceleration Structures

- Mostly used: kd-trees or BVHs
- Must be generated in a preprocess:  $O(n)$ ,  $O(n \log n)$ ,  $O(n^2)$ , ...
- but then traversal is usually  $O(\log n)$
- Performance very much depends on quality of hierarchy
  - good choice of splitting plane !
  - SAH delivers good results

# Acceleration Structures

- Tomorrow:  
Parallel ray traversal on CPUs and GPUs → Kai Selgrad